Avalon® Interface Specifications

Updated for Intel® Quartus® Prime Design Suite: **20.1**

Contents

1. Introduction to the Avalon® Interface Specifications

Avalon® interfaces simplify system design by allowing you to easily connect components in Intel® FPGA. The Avalon interface family defines interfaces appropriate for streaming high-speed data, reading and writing registers and memory, and controlling off-chip devices. Components available in Platform Designer incorporate these standard interfaces. Additionally, you can incorporate Avalon interfaces in custom components, enhancing the interoperability of designs.

This specification defines all the Avalon interfaces. After reading this specification, you should understand which interfaces are appropriate for your components and which signal roles to use for particular behaviors. This specification defines the following seven interfaces:

- Avalon Streaming Interface (Avalon-ST)—an interface that supports the unidirectional flow of data, including multiplexed streams, packets, and DSP data.

- Avalon Memory Mapped Interface (Avalon-MM)—an address-based read/write interface typical of Host-Agent connections.

- Avalon Conduit Interface— an interface type that accommodates individual signals or groups of signals that do not fit into any of the other Avalon types. You can connect conduit interfaces inside a Platform Designer system. Alternatively, you can export them to connect to other modules in the design or to FPGA pins.

- Avalon Tri-State Conduit Interface (Avalon-TC) —an interface to support connections to off-chip peripherals. Multiple peripherals can share pins through signal multiplexing, reducing the pin count of the FPGA and the number of traces on the PCB.

- Avalon Interrupt Interface—an interface that allows components to signal events to other components.

- Avalon Clock Interface—an interface that drives or receives clocks.

- Avalon Reset Interface—an interface that provides reset connectivity.

A single component can include any number of these interfaces and can also include multiple instances of the same interface type.

Note: Avalon interfaces are an open standard. No license or royalty is required to develop and sell products that use or are based on Avalon interfaces.

Related Information

- Introduction to Intel FPGA IP Cores
  Provides general information about all Intel FPGA IP cores, including parameterizing, generating, upgrading, and simulating IP cores.

- Generating a Combined Simulator Setup Script
  Create simulation scripts that do not require manual updates for software or IP version upgrades.

- **Project Management Best Practices**
  Guidelines for efficient management and portability of your project and IP files.

1.1. Avalon Properties and Parameters

Avalon interfaces describe their behavior with properties. The specification for each interface type defines all the interface properties and default values. For example, the `maxChannel` property of Avalon-ST interfaces allows you to specify the number of channels supported by the interface. The `clockRate` property of the Avalon Clock interface provides the frequency of a clock signal.

1.2. Signal Roles

Each Avalon interface defines signal roles and their behavior. Many signal roles are optional. You have the flexibility to select only the signal roles necessary to implement the required functionality. For example, the Avalon-MM interface includes optional `beginbursttransfer` and `burstcount` signal roles for components that support bursting. The Avalon-ST interface includes the optional `startofpacket` and `endofpacket` signal roles for interfaces that support packets.

Except for Avalon Conduit interfaces, each interface may include only one signal of each signal role. Many signal roles allow active-low signals. Active-high signals are generally used in this document.

1.3. Interface Timing

Subsequent chapters of this document include timing information that describes transfers for individual interface types. There is no guaranteed performance for any of these interfaces. Actual performance depends on many factors, including component design and system implementation.

Most Avalon interfaces must not be edge sensitive to signals other than the clock and reset. Other signals may transition multiple times before they stabilize. The exact timing of signals between clock edges varies depending upon the characteristics of the selected Intel FPGA. This specification does not specify electrical characteristics. Refer to the appropriate device documentation for electrical specifications.

1.4. Example: Avalon Interfaces in System Designs

In this example the Ethernet Controller includes six different interface types:

- Avalon-MM
- Avalon-ST
- Avalon Conduit
- Avalon-TC
- Avalon Interrupt
- Avalon Clock.

The Nios® II processor accesses the control and status registers of on-chip components through an Avalon-MM interface. The scatter gather DMAs send and receive data through Avalon-ST interfaces. Four components include interrupt

interfaces serviced by software running on the Nios II processor. A PLL accepts a clock via an Avalon Clock Sink interface and provides two clock sources. Two components include Avalon-TC interfaces to access off-chip memories. Finally, the DDR3 controller accesses external DDR3 memory through an Avalon Conduit interface.

Figure 1. **Avalon Interfaces in a System Design with Scatter Gather DMA Controller and Nios II Processor**

In the following figure, an external processor accesses the control and status registers of on-chip components via an external bus bridge with an Avalon-MM interface. The PCI Express Root Port controls devices on the printed circuit board and the other components of the FPGA by driving an on-chip PCI Express Endpoint with an Avalon-MM host interface. An external processor handles interrupts from five components. A PLL accepts a reference clock via a Avalon Clock sink interface and provides two clock

sources. The flash and SRAM memories share FPGA pins through an Avalon-TC interface. Finally, an SDRAM controller accesses an external SDRAM memory through an Avalon Conduit interface.

Figure 2.   Avalon Interfaces in a System Design with PCI Express Endpoint and External Processor

intel.
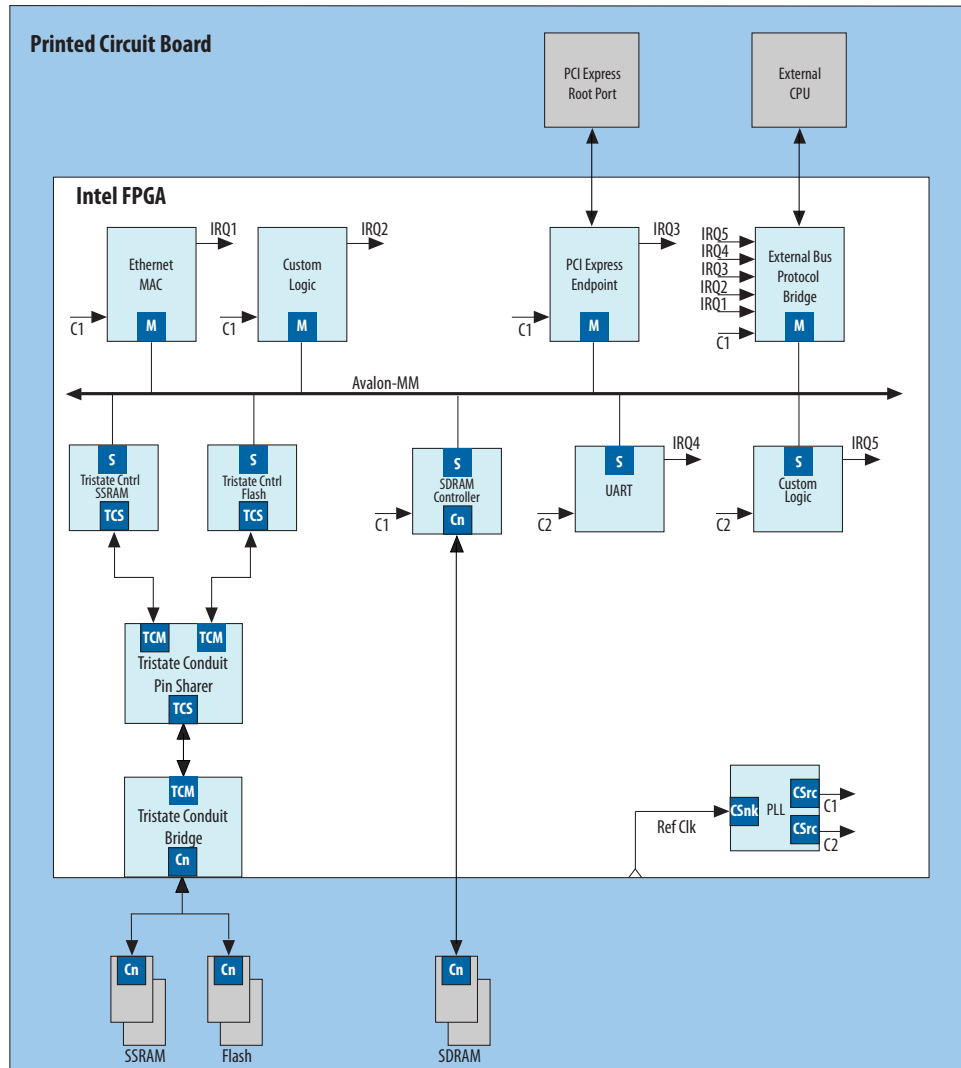
2. Avalon Clock and Reset Interfaces

Avalon Clock interfaces define the clock or clocks used by a component. Components can have clock inputs, clock outputs, or both. A phase locked loop (PLL) is an example of a component that has both a clock input and clock outputs.

The following figure is a simplified illustration showing the most important inputs and outputs of a PLL component.

Figure 3.   PLL Core Clock Outputs and Inputs
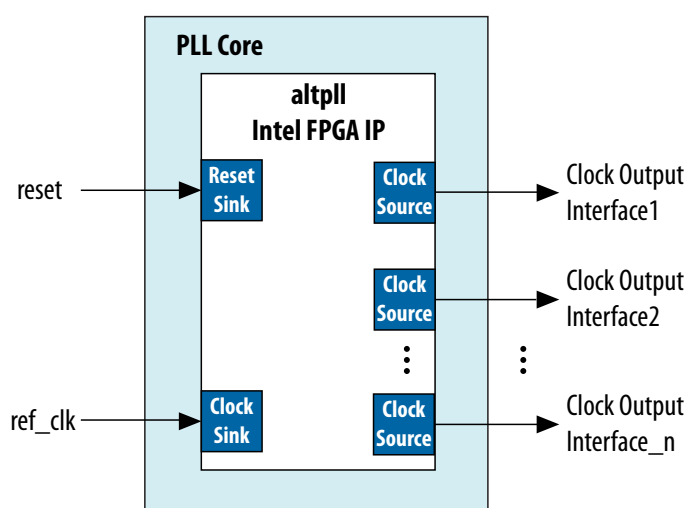
2.1. Avalon Clock Sink Signal Roles

A clock sink provides a timing reference for other interfaces and internal logic.

Table 1.   Clock Sink Signal Roles

Signal Role	Width	Direction	Required	Description
clk	1	Input	Yes	A clock signal. Provides synchronization for internal logic and for other interfaces.

2.2. Clock Sink Properties

Table 2.   Clock Sink Properties

Name	Default Value	Legal Values	Description
clockRate	0	0–2^{32}–1	Indicates the frequency in Hz of the clock sink interface. If 0, the clock rate allows any frequency. If non-zero, Platform Designer issues a warning if the connected clock source is not the specified frequency.

2.3. Associated Clock Interfaces

All synchronous interfaces have an `associatedClock` property that specifies which clock source on the component is used as a synchronization reference for the interface. This property is illustrated in the following figure.

Figure 4.   associatedClock Property

2.4. Avalon Clock Source Signal Roles

An Avalon Clock source interface drives a clock signal out of a component.

Table 3.   Clock Source Signal Roles

Signal Role	Width	Direction	Required	Description
clk	1	Output	Yes	An output clock signal.

2.5. Clock Source Properties

Table 4.   Clock Source Properties

Name	Default Value	Legal Values	Description
associatedDirectClock	N/A	an input clock name	The name of the clock input that directly drives this clock output, if any.
clockRate	0	0–2^{32}–1	Indicates the frequency in Hz at which the clock output is driven.
clockRateKnown	false	true, false	Indicates whether or not the clock frequency is known. If the clock frequency is known, you can customize other components in the system.

Send Feedback

2.6. Reset Sink

Table 5.   Reset Input Signal Roles

The `reset_req` signal is an optional signal that you can use to prevent memory content corruption by performing reset handshake prior to an asynchronous reset assertion.

Signal Role	Width	Direction	Required	Description
reset, reset_n	1	Input	Yes	Resets the internal logic of an interface or component to a user-defined state. The synchronous properties of the reset are defined by the `synchronousEdges` parameter.
reset_req	1	input	No	Early indication of reset signal. This signal acts as a least a one-cycle warning of pending reset for ROM primitives. Use `reset_req` to disable the clock enable or mask the address bus of an on-chip memory, to prevent the address from transitioning when an asynchronous reset input is asserted.

2.7. Reset Sink Interface Properties

Table 6.   Reset Input Signal Roles

Name	Default Value	Legal Values	Description
associatedClock	N/A	a clock name	The name of a clock to which this interface is synchronized. Required if the value of `synchronousEdges` is `DEASSERT` or `BOTH`.
synchronous-Edges	DEASSERT	NONE DEASSERT BOTH	Indicates the type of synchronization the reset input requires. The following values are defined: • `NONE`–no synchronization is required because the component includes logic for internal synchronization of the reset signal. • `DEASSERT`–the reset assertion is asynchronous and deassertion is synchronous. `BOTH`–reset assertion and deassertion are synchronous.

2.8. Associated Reset Interfaces

All synchronous interfaces have an `associatedReset` property that specifies which reset signal resets the interface logic.

2.9. Reset Source

Table 7.   Reset Output Signal Roles

The `reset_req` signal is an optional signal that you can use to prevent memory content corruption by performing reset handshake prior to an asynchronous reset assertion.

Signal Role	Width	Direction	Required	Description
reset reset_n	1	Output	Yes	Resets the internal logic of an interface or component to a user-defined state.
reset_req	1	Output	Optional	Enables reset request generation, which is an early signal that is asserted before reset assertion. Once asserted, this cannot be deasserted until the reset is completed.

2.10. Reset Source Interface Properties

Table 8.   Reset Interface Properties

Name	Default Value	Legal Values	Description
`associatedClock`	N/A	a clock name	The name of a clock to which this interface synchronized. Required if the value of `synchronousEdges` is `DEASSERT` or `BOTH`.
`associatedDirectReset`	N/A	a reset name	The name of the reset input that directly drives this reset source through a one-to-one link.
`associatedResetSinks`	N/A	a reset name	Specifies reset inputs that cause a reset source to assert reset. For example, a reset synchronizer that performs an `OR` operation with multiple reset inputs to generate a reset output.
`synchronousEdges`	`DEASSERT`	NONE DEASSERT BOTH	Indicates the reset output's synchronization. The following values are defined: • `NONE`–The reset interface is asynchronous. • `DEASSERT`–the reset assertion is asynchronous and deassertion is synchronous. • `BOTH`–reset assertion and deassertion are synchronous.

intel

3. Avalon Memory-Mapped Interfaces

3.1. Introduction to Avalon Memory-Mapped Interfaces

You can use Avalon Memory-Mapped (Avalon-MM) interfaces to implement read and write interfaces for Host and Agent components. The following are examples of components that typically include memory-mapped interfaces:

- Microprocessors
- Memories
- UARTs
- DMAs
- Timers

Avalon-MM interfaces range from simple to complex. For example, SRAM interfaces that have fixed-cycle read and write transfers have simple Avalon-MM interfaces. Pipelined interfaces capable of burst transfers are complex.

**ISO
9001:2015
Registered**

Figure 5.   Focus on Avalon-MM Agent Transfers

The following figure shows a typical system, highlighting the Avalon-MM agent interface connection to the interconnect fabric.
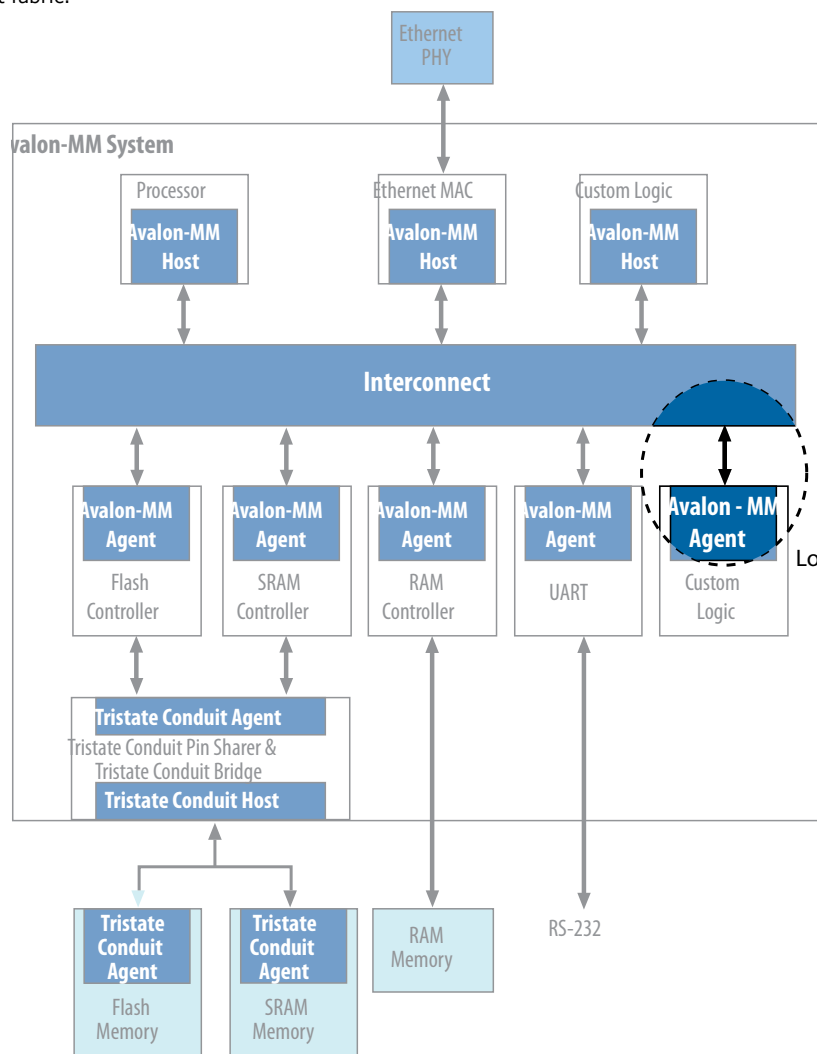
Avalon-MM components typically include only the signals required for the component logic.

Figure 6.   Example Agent Component

The 16-bit general-purpose I/O peripheral shown in the following figure only responds to write requests. This component includes only the Agent signals required for write transfers.

Each signal in an Avalon-MM agent corresponds to exactly one Avalon-MM signal role. An Avalon-MM interface can use only one instance of each signal role.

3.2. Avalon Memory Mapped Interface Signal Roles

Signal roles define the signal types that Avalon memory mapped host and agent ports allow.

This specification does not require all signals to exist in an Avalon memory mapped interface. There is no one signal that is always required. The minimum requirements for an Avalon memory mapped interface are `readdata` for a read-only interface, or `writedata` and `write` for a write-only interface.
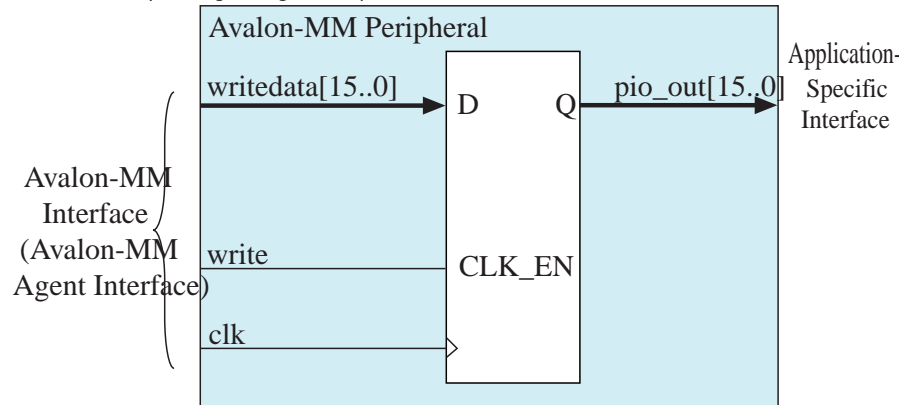
The following table lists signal roles for the Avalon memory mapped interface:

Table 9.   Avalon Memory Mapped Signal Roles

Some Avalon memory mapped signals can be active high or active low. When active low, the signal name ends with _n.

Signal Role	Width	Direction	Required	Description
Fundamental Signals				
address	1 - 64	Host → Agent	No	Hosts: By default, the `address` signal represents a byte address. The value of the address must align to the data width. To write to specific bytes within a data word, the host must use the `byteenable` signal. Refer to the `addressUnits` interface property for word addressing. Agents: By default, the interconnect translates the byte address into a word address in the agent's address space. From the perspective of the agent, each agent access is for a word of data. For example, `address = 0` selects the first word of the agent. `address = 1` selects the second word of the agent. Refer to the `addressUnits` interface property for byte addressing.
byteenable byteenable_n	2, 4, 8, 16, 32, 64, 128	Host → Agent	No	Enables one or more specific byte lanes during transfers on interfaces of width greater than 8 bits. Each bit in `byteenable` corresponds to a byte in `writedata` and `readdata`. The host bit <n> of `byteenable` indicates whether byte <n> is being

continued...

Avalon® Interface Specifications

💬 Send Feedback

Signal Role	Width	Direction	Required	Description
				written to. During writes, `byteenables` specify which bytes are being written to. Other bytes should be ignored by the agent. During reads, `byteenables` indicate which bytes the host is reading. Agents that simply return `readdata` with no side effects are free to ignore `byteenables` during reads. If an interface does not have a `byteenable` signal, the transfer proceeds as if all `byteenables` are asserted. When more than one bit of the `byteenable` signal is asserted, all asserted lanes are adjacent.
debugaccess	1	Host → Agent	No	When asserted, allows the Nios II processor to write on-chip memories configured as ROMs.
read read_n	1	Host → Agent	No	Asserted to indicate a `read` transfer. If present, `readdata` is required.
readdata	8, 16, 32, 64, 128, 256, 512, 1024	Agent → Host	No	The `readdata` driven from the agent to the host in response to a `read` transfer. Required for interfaces that support reads.
response [1:0]	2	Agent → Host	No	The `response` signal is an optional signal that carries the response status. *Note:* Because the signal is shared, an interface cannot issue or accept a write response and a read response in the same clock cycle. • `00: OKAY`—Successful response for a transaction. • `01: RESERVED`—Encoding is reserved. • `10: SLVERR`—Error from an endpoint agent. Indicates an unsuccessful transaction. • `11: DECODEERROR`—Indicates attempted access to an undefined location. For read responses: • One response is sent with each `readdata`. A read burst length of `N` results in `N` responses. Fewer responses are not valid, even in the event of an error. The response signal value may be different for each `readdata` in the burst. • The interface must have read control signals. Pipeline support is possible with the `readdatavalid` signal. • On read errors, the corresponding `readdata` is "don't care". For write responses: • One write response must be sent for each write command. A write burst results in only one response, which must be sent after the final write transfer in the burst is accepted. • If `writeresponsevalid` is present, all write commands must be completed with write responses.
write write_n	1	Host → Agent	No	Asserted to indicate a `write` transfer. If present, `writedata` is required.
writedata	8, 16, 32, 64, 128, 256, 512, 1024	Host → Agent	No	Data for write transfers. The width must be the same as the width of `readdata` if both are present. Required for interfaces that support writes.
Wait-State Signals				

continued...

Signal Role	Width	Direction	Required	Description
lock	1	Host → Agent	No	lock ensures that once a host wins arbitration, the winning host maintains access to the agent for multiple transactions. Lock asserts coincident with the first read or write of a locked sequence of transactions. Lock deasserts on the final transaction of a locked sequence of transactions. lock assertion does not guarantee that arbitration is won. After the lock-asserting host has been granted, that host retains grant until lock is deasserted. A host equipped with lock cannot be a burst host. Arbitration priority values for lock-equipped hosts are ignored. lock is particularly useful for read-modify-write (RMW) operations. The typical read-modify-write operation includes the following steps: 1. Host A asserts lock and reads 32-bit data that has multiple bit fields. 2. Host A deasserts lock, changes one bit field, and writes the 32-bit data back. lock prevents host B from performing a write between Host A's read and write.
waitrequest waitrequest_n	1	Agent → Host	No	An agent asserts waitrequest when unable to respond to a read or write request. Forces the host to wait until the interconnect is ready to proceed with the transfer. At the start of all transfers, a host initiates the transfer and waits until waitrequest is deasserted. A host must make no assumption about the assertion state of waitrequest when the host is idle: waitrequest may be high or low, depending on system properties. When waitrequest is asserted, host control signals to the agent must remain constant except for beginbursttransfer. For a timing diagram illustrating the beginbursttransfer signal, refer to the figure in *Read Bursts*. An Avalon memory mapped agent may assert waitrequest during idle cycles. An Avalon memory mapped host may initiate a transaction when waitrequest is asserted and wait for that signal to be deasserted. To avoid system lockup, an agent device should assert waitrequest when in reset.
Pipeline Signals				
readdatavalid readdatavalid_n	1	Agent → Host	No	Used for variable-latency, pipelined read transfers. When asserted, indicates that the readdata signal contains valid data. For a read burst with burstcount value *<n>*, the readdatavalid signal must be asserted *<n>* times, once for each readdata item. There must be at least one cycle of latency between acceptance of the read and assertion of readdatavalid. For a timing diagram illustrating the readdatavalid signal, refer to *Pipelined Read Transfer with Variable Latency*. An agent may assert readdatavalid to transfer data to the host independently of whether the agent is stalling a new command with waitrequest. Required if the host supports pipelined reads. Bursting hosts with read functionality must include the readdatavalid signal.
writeresponsevalid	1	Agent → Host	No	An optional signal. If present, the interface issues write responses for write commands. When asserted, the value on the response signal is a valid write response. Writeresponsevalid is only asserted one clock cycle or more after the write command is accepted. There is at least a one clock cycle latency from command acceptance to assertion of writeresponsevalid.

continued...

Send Feedback

Signal Role	Width	Direction	Required	Description
				A write command is considered accepted when the last beat of the burst is issued to the agent and `waitrequest` is low. `writeresponsevalid` can be asserted one or more clock cycles after the last beat of the burst has been issued.
Burst Signals				
`burstcount`	1 – 11	Host → Agent	No	Used by bursting hosts to indicate the number of transfers in each burst. The value of the maximum `burstcount` parameter must be a power of 2. A burstcount interface of width <n> can encode a max burst of size $2^{(<n>-1)}$. For example, a 4-bit `burstcount` signal can support a maximum burst count of 8. The minimum `burstcount` is 1. The `constantBurstBehavior` property controls the timing of the `burstcount` signal. Bursting hosts with read functionality must include the `readdatavalid` signal._For bursting hosts and agents using byte addresses, the following restriction applies to the width of the address:_`<address_w> >=` `<burstcount_w> +``log₂(<symbols_per_word_of_interface>)`For bursting hosts and agents using word addresses, the $\log_2$ term above is omitted.
`beginbursttransfer`	1	Interconnect → Agent	No	Asserted for the first cycle of a burst to indicate when a burst transfer is starting. This signal is deasserted after one cycle regardless of the value of `waitrequest`. For a timing diagram illustrating `beginbursttransfer`, refer to the figure in _Read Bursts_._`beginbursttransfer` is optional. An agent can always internally calculate the start of the next write burst transaction by counting data transfers._**Warning:** _do not_ use this signal. This signal exists to support legacy memory controllers.

3.3. Interface Properties

Table 10.   Avalon-MM Interface Properties

Name	Default Value	Legal Values	Description
`addressUnits`	Host - symbols Agent - words	words, symbols	Specifies the unit for addresses. A symbol is typically a byte.Refer to the definition of `address` in the _Avalon Memory-Mapped Interface Signal Types_ table for the typical use of this property.
`alwaysBurstMaxBurst`	false	true, false	When true, indicates that the host always issues the maximum-length burst. The maximum burst length is $2^{burstcount\_width\ -\ 1}$. This parameter has no effect for Avalon-MM agent interfaces.
`burstcountUnits`	words	words, symbols	This property specifies the units for the burstcount signal. For symbols, the `burstcount` value is interpreted as the number of symbols (bytes) in the burst. For words, the `burstcount` value is interpreted as the number of word transfers in the burst.
`burstOnBurstBoundariesOnly`	false	true, false	If true, burst transfers presented to this interface begin at addresses which are multiples of the maximum burst size.
			continued...

Name	Default Value	Legal Values	Description
`constantBurstBehavior`	Host -false Agent -false	true, false	Hosts: When true, declares that the host holds address and burstcount constant throughout a burst transaction. When false (default), declares that the host holds address and burstcount constant only for the first beat of a burst. Agents: When true, declares that the agent expects address and burstcount to be held constant throughout a burst. When false (default), declares that the agent samples address and burstcount only on the first beat of a burst.
`holdTime`(1)	0	0 – 1000 cycles	Specifies time in `timingUnits` between the deassertion of `write` and the deassertion of `address` and `data`. (Only applies to write transactions.)
`linewrapBursts`	false	true, false	Some memory devices implement a wrapping burst instead of an incrementing burst. When a wrapping burst reaches a burst boundary, the address wraps back to the previous burst boundary. Only the low-order bits are required for address counting. For example, a wrapping burst to address 0xC with burst boundaries every 32 bytes across a 32-bit interface writes to the following addresses: • 0xC • 0x10 • 0x14 • 0x18 • 0x1C • 0x0 • 0x4 • 0x8
`maximumPendingReadTransactions` (1)	1(2)	1 – 64	Agents: This parameter is the maximum number of pending reads that the agent can queue. The value must be non-zero for any agent with the `readdatavalid` signal. Refer to *Pipelined Read Transfer with Variable Latency* for a timing diagram that illustrates this property and for additional information about using `waitrequest` and `readdatavalid` with multiple outstanding reads. Hosts: This property is the maximum number of outstanding read transactions that the host can generate. *Note:* Do not set this parameter to 0. (For backwards compatibility, the software supports a parameter setting of 0. However, you should not use this setting in new designs).
`maximumPendingWriteTransactions`	0	1 – 64	The maximum number of pending non-posted writes that a agent can accept or a host can issue. A agent asserts `waitrequest` once the interconnect reaches this limit, and the host stops issuing commands. The default value is 0, which allows unlimited pending write transactions for a host that supports write responses. A agent that supports write responses must set this to a non-zero value.
`minimumResponseLatency`	1		For interfaces that support `readdatavalid` or `writeresponsevalid`, specifies the minimum number of cycles between a read or write command and the response to the command.

continued...

Send Feedback

Name	Default Value	Legal Values	Description
readLatency(1)	0	0 – 63	Read latency for fixed-latency Avalon-MM agents. For a timing diagram that uses a fixed latency read, refer to *Pipelined Read Transfers with Fixed Latency*. Avalon-MM agents that are fixed latency must provide a value for this interface property. Avalon-MM agents that are variable latency use the readdatavalid signal to specify valid data.
readWaitTime(1)	1	0 – 1000 cycles	For interfaces that do not use the waitrequest signal. readWaitTime indicates the timing in timingUnits before the agents accepts a read command. The timing is as if the agent asserted waitrequest for readWaitTime cycles.
setupTime(1)	0	0 – 1000 cycles	Specifies time in timingUnits between the assertion of address and data and assertion of read or write.
timingUnits(1)	cycles	cycles, nanoseconds	Specifies the units for setupTime, holdTime, writeWaitTime and readWaitTime. Use cycles for synchronous devices and nanoseconds for asynchronous devices. Almost all Avalon-MM agent devices are synchronous. An Avalon-MM component that bridges from an Avalon-MM agent interface to an off-chip device may be asynchronous. That off-chip device might have a fixed settling time for bus turnaround.
waitrequestAllowance	0		Specifies the number of transfers that can be issued or accepted after waitrequest is asserted. When the waitrequestAllowance is 0, the write, read and waitrequest signals maintain their existing behavior as described in the *Avalon-MM Signal Roles* table. When the waitrequestAllowance is greater than 0, every clock cycle on which write or read is asserted counts as a command transfer. Once waitrequest is asserted, only waitrequestAllowance more command transfers are legal while waitrequest remains asserted. After the waitrequestAllowance is reached, write and read must remain deasserted for as long as waitrequest is asserted. Once waitrequest deasserts, transfers may resume at any time without restrictions until waitrequest asserts again. At this time, waitrequestAllowance more transfers may complete while waitrequest remains asserted.
writeWaitTime(1)	0	0 – 1000 Cycles	For interfaces that do not use the waitrequest signal, writeWaitTime specifies the timing in timingUnits before a agent accepts a write. The timing is as if the agent asserted waitrequest for writeWaitTime cycles or nanoseconds. For a timing diagram that illustrates the use of writeWaitTime, refer to *Read and Write Transfers with Fixed Wait-States*.
Interface Relationship Properties			
associatedClock	N/A	N/A	Name of the clock interface to which this Avalon-MM interface is synchronous.

Name	Default Value	Legal Values	Description
`associatedReset`	N/A	N/A	Name of the reset interface which resets the logic on this Avalon-MM interface.
`bridgesToHost`	0	Avalon-MM Host name on the same component	An Avalon-MM bridge consists of a agent and a host, and has the property that an access to the agent requesting a byte or bytes causes the same byte or bytes to be requested by the host. The Avalon-MM Pipeline Bridge in the Platform Designer component library implements this functionality.

Notes:
1. Although this property characterizes a agent device, hosts can declare this property to enable direct connections between matching host and agent interfaces.
2. If a agent interface accepts more read transfers than allowed, the interconnect pending read FIFO may overflow with unpredictable results. The agent may lose `readdata` or route `readdata` to the wrong host interface. Or, the system may lock up. The agent interface must assert `waitrequest` to prevent this overflow.

Related Information

- Avalon Memory Mapped Interface Signal Roles on page 14
- Read and Write Responses on page 34
- Pipelined Read Transfer with Variable Latency on page 28
- Pipelined Read Transfers with Fixed Latency on page 29
- Read and Write Responses
  In *Platform Designer User Guide: Intel Quartus® Prime Pro Edition*

3.4. Timing

The Avalon-MM interface is synchronous. Each Avalon-MM interface is synchronized to an associated clock interface. Signals may be combinational if they are driven from the outputs of registers that are synchronous to the clock signal. This specification does not dictate how or when signals transition between clock edges. Timing diagrams are devoid of fine-grained timing information.

3.5. Transfers

This section defines two basic concepts before introducing the transfer types:

- Transfer—A transfer is a read or write operation of a word or one or more symbol of data. Transfers occur between an Avalon-MM interface and the interconnect. Transfers take one or more clock cycles to complete.

  Both hosts and agents are part of a transfer. The Avalon-MM host initiates the transfer and the Avalon-MM agent responds.

- Host-Agent pair—This term refers to the host interface and agent interface involved in a transfer. During a transfer, the host interface control and data signals pass through the interconnect fabric and interact with the agent interface.

3.5.1. Typical Read and Write Transfers

This section describes a typical Avalon-MM interface that supports read and write transfers with agent-controlled `waitrequest`. The agent can stall the interconnect for as many cycles as required by asserting the `waitrequest` signal. If a agent uses `waitrequest` for either read or write transfers, the agent must use `waitrequest` for both.

A agent typically receives `address`, `byteenable`, `read` or `write`, and `writedata` after the rising edge of the clock. A agent asserts `waitrequest` before the rising clock edge to hold off transfers. When the agent asserts `waitrequest`, the transfer is delayed. While `waitrequest` is asserted, the address and other control signals are held constant. Transfers complete on the rising edge of the first `clk` after the agent interface deasserts `waitrequest`.

There is no limit on how long a agent interface can stall. Therefore, you must ensure that a agent interface does not assert `waitrequest` indefinitely. The following figure shows `read` and `write` transfers using `waitrequest`.

Note: `waitrequest` can be decoupled from the `read` and `write` request signals. `waitrequest` may be asserted during idle cycles. An Avalon-MM host may initiate a transaction when `waitrequest` is asserted and wait for that signal to be deasserted. Decoupling `waitrequest` from `read` and `write` requests may improve system timing. Decoupling eliminates a combinational loop including the `read`, `write`, and `waitrequest` signals. If even more decoupling is required, use the `waitrequestAllowance` property. `waitrequestAllowance` is available starting with the Quartus® Prime Pro v17.1 Stratix® 10 ES Editions release.

Figure 7.   Read and Write Transfers with Waitrequest
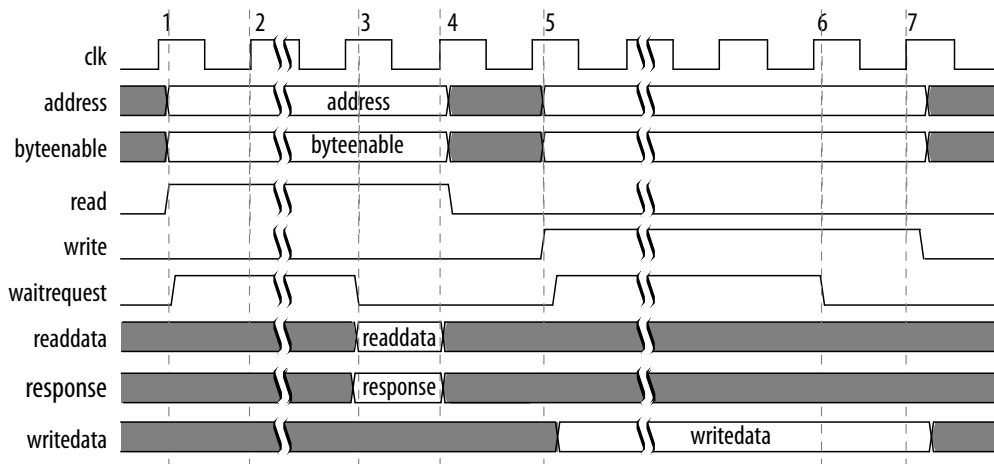
The numbers in this timing diagram, mark the following transitions:

1. `address`, `byteenable`, and `read` are asserted after the rising edge of `clk`. The agent asserts `waitrequest`, stalling the transfer.

2. `waitrequest` is sampled. Because `waitrequest` is asserted, the cycle becomes a wait-state. `address`, `read`, `write`, and `byteenable` remain constant.

3. The agent deasserts `waitrequest` after the rising edge of `clk`. The agent asserts `readdata` and `response`.

4. The host samples `readdata`, `response` and deasserted `waitrequest` completing the transfer.

5. `address`, `writedata`, `byteenable`, and `write` signals are asserted after the rising edge of `clk`. The agent asserts `waitrequest` stalling the transfer.

6. The agent deasserts `waitrequest` after the rising edge of `clk`.

7. The agent captures write data ending the transfer.

3.5.2. Transfers Using the waitrequestAllowance Property

The `waitrequestAllowance` property specifies the number of transfers an Avalon-MM host can issue or an Avalon-MM agent must accept after the `waitrequest` signal is asserted. `waitrequestAllowance` is available starting with the Intel Quartus Prime 17.1 software release.

The default value of `waitrequestAllowance` is 0, which corresponds to the behavior described in *Typical Read and Write Transfers*, where `waitrequest` assertion stops the current transfer from being issued or accepted.

An Avalon-MM agent with a `waitrequestAllowance` greater than 0 would typically assert `waitrequest` when its internal buffer can only accept `waitrequestAllowance` more entries before becoming full. Avalon-MM hosts with a `waitrequestAllowance` greater than 0 have `waitrequestAllowance` additional cycles to stop sending transfers, which allows more pipelining in the host logic. The host must deassert the `read` or `write` signal when the `waitrequestallowance` has been spent.

Values of `waitrequestAllowance` greater than 0 support high-speed design where immediate forms of backpressure may result in a drop in the maximum operating frequency (F_{MAX}) often due to combinatorial logic in the control path. An Avalon-MM agent must support all possible transfer timings that are legal for its `waitrequestAllowance` value. For example, a agent with `waitrequestAllowance = 2` must be able to accept any of the host transfer waveforms shown in the following examples.

Related Information

3.5.2.1. waitrequestAllowance Equals Two

The following timing diagram illustrates timing for an Avalon-MM host that has two clock cycles to start and stop sending transfers after the Avalon-MM agent deasserts or asserts `waitrequest`, respectively.

Figure 8. **Host write: waitrequestAllowance Equals Two Clock Cycles**

The markers in this figure mark the following events:

1. The Avalon-MM> host drives `write` and `data`.

2. The Avalon-MM> agent asserts `waitrequest`. Because the `waitrequestAllowance` is 2, the host is able to complete the 2 additional data transfers.

3. The host deasserts `write` as required because the agent is asserting `waitrequest` for a third cycle.

4. The Avalon-MM> host drives `write` and `data`. The agent is not asserting `waitrequest`. The writes complete.

5. The Avalon host drives `write` and `data` even though the agent is asserting `waitrequest`. Because the `waitrequestAllowance` is 2 cycles, the write completes.

6. The Avalon host drives `write` and `data`. The agent is not asserting `waitrequest`. The write completes.

3.5.2.2. waitrequestAllowance Equals One

The following timing diagram illustrates timing for an Avalon-MM host that has one clock cycle to start and stop sending transfers after the Avalon-MM agent deasserts or asserts `waitrequest`, respectively:

**Figure 9.   Host Write: waitrequestAllowance Equals One Clock Cycle**

The numbers in this figure mark the following events:

1. The Avalon-MM host drives `write` and `data`.

2. The Avalon-MM agent asserts `waitrequest`. Because the `waitrequestAllowance` is 1, the host can complete the write.

3. The host deasserts `write` because the agent is asserting `waitrequest` for a second cycle.

4. The Avalon-MM host drives `write` and `data`. The agent is not asserting `waitrequest`. The writes complete.

5. The agent asserts `waitrequest`. Because the `waitrequestAllowance` is 1 cycle, the write completes.

6. Avalon-MM host drives `write` and `data`. The agent is not asserting `waitrequest`. The write completes.

7. The Avalon-MM agent asserts `waitrequest`. Because the `waitrequestAllowance` is 1, the host can complete one additional data transfer.

8. The Avalon host drives `write` and `data`. The agent is not asserting `waitrequest`. The write completes.

3.5.2.3. waitrequestAllowance Equals Two - Not Recommended

The following diagram illustrates timing for an Avalon-MM> host that can send two transfers after `waitrequest` is asserted.

This timing is legal, but not recommended. In this example the host counts the number of transactions instead of the number of clock cycles. This approach requires a counter that makes the implementation more complex and may affect timing closure. When the host determines when to drive transactions with the `waitrequest` signal and a constant number of cycles, the host starts or stops transactions based on the registered signals.

Figure 10.   waitrequestAllowance Equals Two Transfers

The numbers in this figure mark the following events:

1. The Avalon-MM> host asserts `write` and drives `data`.

2. The Avalon-MM> agent asserts `waitrequest`.

3. The Avalon-MM> host drives `write` and `data`. Because the `waitrequestAllowance` is 2, the host drives data in 2 consecutive cycles.

4. The Avalon-MM> host deasserts `write` because the host has spent the 2-transfer `waitrequestAllowance`.

5. The Avalon-MM> host issues a write as soon as `waitrequest` is deasserted.

6. The Avalon-MM> host drives `write` and `data`. The agent asserts `waitrequest` for 1 cycle.

7. In response to `waitrequest`, the Avalon-MM> host holds data for 2 cycles.

3.5.2.4. waitrequestAllowance Compatibility for Avalon-MM Host and Agent Interfaces

Avalon-MM hosts and agents that support the `waitrequest` signal support backpressure. Hosts with backpressure can always connect to agents without backpressure. Hosts without backpressure cannot connect to agents with backpressure.

Table 11.   waitrequestAllowance Compatibility for Avalon-MM Hosts and Agents

Host and Agent waitrequestAllowance	Compatibility
host = 0 agent = 0	Follows the same compatibility rules as standard Avalon-MM interfaces.
host = 0 agent > 0	Direct connections are not possible. Simple adaptation is required for the case of a host with a `waitrequest` signal. A connection is impossible if the host does not support the `waitrequest` signal.
host > 0 agent = 0	Direct connections are not possible. Adaptation (buffers) are required when connecting to a agent with a `waitrequest` signal or fixed wait states.
host > 0 agent> 0	No adaptation is required if the host's allowance <= agent's allowance. If the host allowance < agent allowance, pipeline registers may be inserted. For point-to-point connections, you can add the pipeline registers on the command signals or the `waitrequest` signals. Up to <d> register stages can be inserted where <d> is the difference between the allowances. Connecting a host with a higher `waitrequestAllowance` than the agent requires buffering.

3.5.2.5. waitrequestAllowance Error Conditions

Behavior is unpredictable for if an Avalon-MM interface violates the `waitrequest` allowance specification.

- If a host violates the `waitrequestAllowance` = *<n>* specification by sending more than *<n>* transfers, transfers may be dropped or data corruption may occur.

- If a agent advertises a larger `waitrequestAllowance` than is possible, some transfers may be dropped or data corruption may occur.

3.5.3. Read and Write Transfers with Fixed Wait-States

A agent can specify fixed wait-states using the `readWaitTime` and `writeWaitTime` properties. Using fixed wait-states is an alternative to using `waitrequest` to stall a transfer. The address and control signals (`byteenable`, `read`, and `write`) are held constant for the duration of the transfer. Setting `readWaitTime` or `writeWaitTime` to *<n>* is equivalent to asserting `waitrequest` for *<n>* cycles per transfer.

In the following figure, the agent has a `writeWaitTime` = 2 and `readWaitTime` = 1.

Figure 11.   Read and Write Transfer with Fixed Wait-States at the Agent Interface

The numbers in this timing diagram mark the following transitions:

1. The host asserts `address` and `read` on the rising edge of `clk`.

2. The next rising edge of `clk` marks the end of the first and only wait-state cycle. The `readWaitTime` is 1.

3. The agent asserts `readdata` and `response` on the rising edge of `clk`. The read transfer ends.

4. `writedata`, `address`, `byteenable`, and `write` signals are available to the agent.

5. The write transfer ends after 2 wait-state cycles.

Transfers with a single wait-state are commonly used for multicycle off-chip peripherals. The peripheral captures address and control signals on the rising edge of `clk`. The peripheral has one full cycle to return data.

Components with zero wait-states are allowed. However, components with zero wait-states may decrease the achievable frequency. Zero wait-states require the component to generate the response in the same cycle that the request was presented.

3.5.4. Pipelined Transfers

Avalon-MM pipelined read transfers increase the throughput for synchronous agent devices that require several cycles to return data for the first access. Such devices can typically return one data value per cycle for some time thereafter. New pipelined read transfers can start before `readdata` for the previous transfers is returned.

A pipelined read transfer has an address phase and a data phase. A host initiates a transfer by presenting the address during the address phase. A agent fulfills the transfer by delivering the data during the data phase. The address phase for a new transfer (or multiple transfers) can begin before the data phase of a previous transfer completes. The delay is called pipeline latency. The pipeline latency is the duration from the end of the address phase to the beginning of the data phase.

Transfer timing for wait-states and pipeline latency have the following key differences:

- Wait-states—Wait-states determine the length of the address phase. Wait-states limit the maximum throughput of a port. If a agent requires one wait-state to respond to a transfer request, the port requires two clock cycles per transfer.

- Pipeline Latency—Pipeline latency determines the time until data is returned independently of the address phase. A pipelined agent with no wait-states can sustain one transfer per cycle. However, the agent may require several cycles of latency to return the first unit of data.

Wait-states and pipelined reads can be supported concurrently. Pipeline latency can be either fixed or variable.

3.5.4.1. Pipelined Read Transfer with Variable Latency

After capturing address and control signals, an Avalon-MM pipelined agent takes one or more cycles to produce data. A pipelined agent may have multiple pending read transfers at any given time.

Variable-latency pipelined read transfers:

- Require one additional signal, `readdatavalid`, that indicates when read data is valid.

- Include the same set of signals as non-pipelined read transfers.

In variable-latency pipelined read transfers, Agent peripherals that use `readdatavalid` are considered pipelined with variable latency. The `readdata` and `readdatavalid` signals corresponding to a read command can be asserted the cycle after that read command is asserted, at the earliest.

The agent must return `readdata` in the same order that the read commands are accepted. Pipelined agent ports with variable latency must use `waitrequest`. The agent can assert `waitrequest` to stall transfers to maintain an acceptable number of pending transfers. A agent may assert `readdatavalid` to transfer data to the host independently of whether the agent is stalling a new command with `waitrequest`.

Note: The maximum number of pending transfers is a property of the agent interface. The interconnect fabric builds logic to route `readdata` to requesting hosts using this number. The agent interface, not the interconnect fabric, must track the number of pending reads. The agent must assert `waitrequest` to prevent the number of pending reads from exceeding the maximum number. If a agent has `waitrequestAllowance > 0`, the agent must assert `waitrequest` early enough so that the total pending transfers, including those accepted while `waitrequest` is asserted, does not exceed the maximum number of pending transfers specified.

Figure 12.   Pipelined Read Transfers with Variable Latency

The following figure shows several agent read transfers. The agent is pipelined with variable latency. In this figure, the agent can accept a maximum of two pending transfers. The agent uses `waitrequest` to avoid overrunning this maximum.

The numbers in this timing diagram, mark the following transitions:

1. The host asserts `address` and `read`, initiating a read transfer.

2. The agent captures `addr1`.

3. The agent captures `addr2`.

4. The agent asserts `waitrequest` because the agent has already accepted a maximum of two pending reads, causing the third transfer to stall.

5. The agent asserts `data1`, the response to `addr1`. The agent deasserts `waitrequest`.

6. The agent captures `addr3`. The interconnect captures `data1`.

7. The agent captures `addr4`. The interconnect captures `data2`.

8. The agent drives `readdatavalid` and `readdata` in response to the third read transfer.

9. The agent captures `addr5`. The interconnect captures `data3`. The `read` signal is deasserted. The value of `waitrequest` is no longer relevant.

10. The interconnect captures `data4`.

11. The agent drives `data5` and asserts `readdatavalid` completing the data phase for the final pending read transfer.

If the agent cannot handle a write transfer while processing pending read transfers, the agent must assert `waitrequest` and stall the write operation until the pending read transfers have completed. The Avalon-MM specification does not define the value of `readdata` in the event that a agent accepts a write transfer to the same address as a currently pending read transfer.

3.5.4.2. Pipelined Read Transfers with Fixed Latency

The address phase for fixed latency read transfers is identical to the variable latency case. After the address phase, a pipelined with fixed read latency takes a fixed number of clock cycles to return valid `readdata`. The `readLatency` property specifies the number of clock cycles to return valid `readdata`. The interconnect captures `readdata` on the appropriate rising clock edge, ending the data phase.

During the address phase, the can assert `waitrequest` to hold off the transfer. Or, the specifies the `readLatency` for a fixed number of wait states. The address phase ends on the next rising edge of `clk` after wait states, if any.

During the data phase, the drives `readdata` after a fixed latency. For a read latency of *<n>*, the must present valid `readdata` on the *<nth>* rising edge of `clk` after the end of the address phase.

Figure 13.   Pipelined Read Transfer with Fixed Latency of Two Cycles

The following figure shows multiple data transfers between a host and a pipelined . The drives `waitrequest` to stall transfers and has a fixed read latency of 2 cycles.

The numbers in this timing diagram, mark the following transitions:

1. A host initiates a read transfer by asserting `read` and `addr1`.

2. The asserts `waitrequest` to hold off the transfer for one cycle.

3. The captures `addr1` at the rising edge of `clk`. The address phase ends here.

4. The presents valid `readdata` after 2 cycles, ending the transfer.

5. `addr2` and `read` are asserted for a new read transfer.

6. The host initiates a third read transfer during the next cycle, before the data from the prior transfer is returned.

3.5.5. Burst Transfers

A burst executes multiple transfers as a unit, rather than treating every word independently. Bursts may increase throughput for agent ports that achieve greater efficiency when handling multiple words at a time, such as SDRAM. The net effect of bursting is to lock the arbitration for the duration of the burst. A bursting Avalon-MM interface that supports both reads and writes must support both read and write bursts.

Bursting Avalon-MM interfaces include a `burstcount` output signal. If a agent has a `burstcount` input, the agent is burst capable.

The `burstcount` signal behaves as follows:

*   At the start of a burst, `burstcount` presents the number of sequential transfers in the burst.

*   For width *<n>* of `burstcount`, the maximum burst length is $2^{(<n>-1)}$. The minimum legal burst length is one.

To support agent read bursts, a agent must also support:

- Wait states with the `waitrequest` signal.
- Pipelined transfers with variable latency with the `readdatavalid` signal.

At the start of a burst, the agent sees the `address` and a burst length value on `burstcount`. For a burst with an address of *<a>* and a `burstcount` value of *<b>*, the agent must perform *<b>* consecutive transfers starting at address *<a>*. The burst completes after the agent receives (write) or returns (read) the *<b*th*>* word of data. The bursting agent must capture `address` and `burstcount` only once for each burst. The agent logic must infer the address for all but the first transfers in the burst. A agent can also use the input signal `beginbursttransfer`, which the interconnect asserts on the first cycle of each burst.

3.5.5.1. Write Bursts

These rules apply when a write burst begins with `burstcount` greater than one:

- When a `burstcount` of *<n>* is presented at the beginning of the burst, the agent must accept *<n>* successive units of `writedata` to complete the burst. Arbitration between the host-agent pair remains locked until the burst completes. This lock guarantees that no other host can execute transactions on the agent until the write burst completes.

- The agent must only capture `writedata` when `write` asserts. During the burst, the host can deassert `write` indicating that `writedata` is invalid. Deasserting `write` does not terminate the burst. The `write` deassertion delays the burst and no other host can access the agent, reducing the transfer efficiency.

- The agent delays a transfer by asserting `waitrequest` forcing `writedata`, `write`, `burstcount`, and `byteenable` to be held constant.

- The functionality of the `byteenable` signal is the same for bursting and non-bursting agents. For a 32-bit host burst-writing to a 64-bit agent, starting at byte address 4, the first write transfer seen by the agent is at its address 0, with `byteenable = 8'b11110000`. The `byteenables` can change for different words of the burst.

- The `byteenable` signals do not all have to be asserted. A burst host writing partial words can use the `byteenable` signal to identify the data being written.

- Writes with `byteenable` signals being all 0's are simply passed on to the Avalon-MM agent as valid transactions.

- The `constantBurstBehavior` property specifies the behavior of the burst signals.

  — When `constantBurstBehavior` is true for a host, the host holds `address` and `burstcount` stable throughout a burst. When true for a agent, `constantBurstBehavior` declares that the agent expects `address` and `burstcount` to be held stable throughout a burst.

  — When `constantBurstBehavior` is false, the host holds `address` and `burstcount` stable only for the first transaction of a burst. When `constantBurstBehavior` is false, the agent samples `address` and `burstcount` only on the first transaction of a burst.

Figure 14. **Write Burst with constantBurstBehavior Set to False for Host and Agent**

The following figure demonstrates a agent write burst of length 4. In this example, the agent asserts `waitrequest` twice delaying the burst.
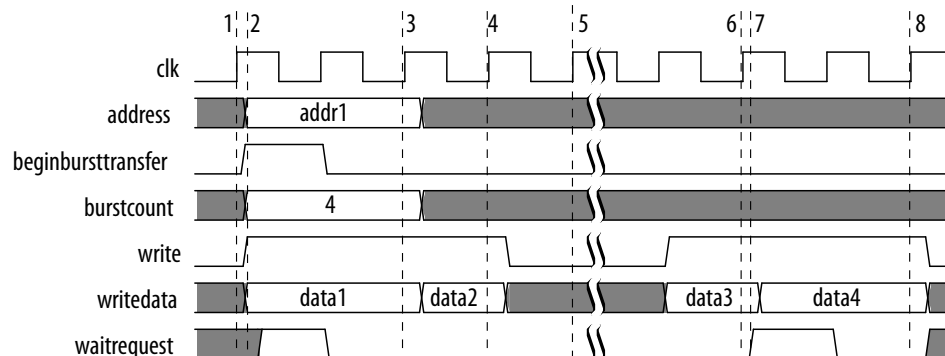
The numbers in this timing diagram mark the following transitions:

1. The host asserts `address`, `burstcount`, `write`, and drives the first unit of `writedata`.

2. The agent immediately asserts `waitrequest`, indicating that the agent is not ready to proceed with the transfer.

3. `waitrequest` is low. The agent captures `addr1`, `burstcount`, and the first unit of `writedata`. On subsequent cycles of the transfer, `address` and `burstcount` are ignored.

4. The agent captures the second unit of data at the rising edge of `clk`.

5. The burst is paused while `write` is deasserted.

6. The agent captures the third unit of data at the rising edge of `clk`.

7. The agent asserts `waitrequest`. In response, all outputs are held constant through another clock cycle.

8. The agent captures the last unit of data on this rising edge of `clk`. The agent write burst ends.

In the figure above, the `beginbursttransfer` signal is asserted for the first clock cycle of a burst and is deasserted on the next clock cycle. Even if the agent asserts `waitrequest`, the `beginbursttransfer` signal is only asserted for the first clock cycle.

Related Information

Interface Properties on page 17

3.5.5.2. Read Bursts

Read bursts are similar to pipelined read transfers with variable latency. A read burst has distinct address and data phases. `readdatavalid` indicates when the agent is presenting valid `readdata`. Unlike pipelined read transfers, a single read burst address results in multiple data transfers.

These rules apply to read bursts:

- When a host connects directly to a agent, a `burstcount` of *<n>* means the agent must return *<n>* words of `readdata` to complete the burst. For cases where interconnect links the host and agent pair, the interconnect may suppress read commands sent from the host to the agent. For example, if the host sends a read command with a `byteenable` value of 0, the interconnect may suppress the read. As a result, the agent does not respond to the read command.

- The agent presents each word by providing `readdata` and asserting `readdatavalid` for a cycle. Deassertion of `readdatavalid` delays but does not terminate the burst data phase.

- For reads with a `burstcount > 1`, Intel recommends asserting all `byteenables`.

Note: Intel recommends that burst capable agents not have read side effects. (This specification does not guarantee how many bytes a host reads from the agent in order to satisfy a request.)

Figure 15.   Read Burst

The following figure illustrates a system with two bursting hosts accessing a agent. Note that Host B can drive a read request before the data has returned for Host A.
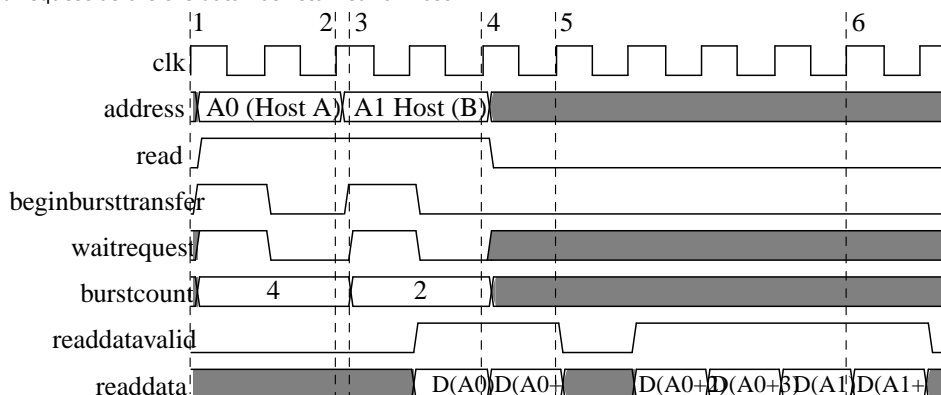
The numbers in this timing diagram, mark the following transitions:

1. Host A asserts `address` (A0), `burstcount`, and `read` after the rising edge of `clk`. The agent asserts `waitrequest`, causing all inputs except `beginbursttransfer` to be held constant through another clock cycle.

2. The agent captures A0 and `burstcount` at this rising edge of `clk`. A new transfer could start on the next cycle.

3. Host B drives `address` (A1), `burstcount`, and `read`. The agent asserts `waitrequest`, causing all inputs except `beginbursttransfer` to be held constant. The agent could have returned read data from the first read request at this time, at the earliest.

4. The agent presents valid `readdata` and asserts `readdatavalid`, transferring the first word of data for host A.

5. The second word for host A is transferred. The agent deasserts `readdatavalid` pausing the read burst. The agent port can keep `readdatavalid` deasserted for an arbitrary number of clock cycles.

6. The first word for host B is returned.

3.5.5.3. Line–Wrapped Bursts

Processors with instruction caches gain efficiency by using line-wrapped bursts. When a processor requests data that is not in the cache, the cache controller must refill the entire cache line. For a processor with a cache line size of 64 bytes, a cache miss causes 64 bytes to be read from memory. If the processor reads from address 0xC when the cache miss occurred, then an inefficient cache controller could issue a burst at address 0, resulting in data from read addresses 0x0, 0x4, 0x8, 0xC, 0x10, 0x14, 0x18, . . . 0x3C. The requested data is not available until the fourth read. With line-wrapping bursts, the address order is 0xC, 0x10, 0x14, 0x18, . . . 0x3C, 0x0, 0x4, and 0x8. The requested data is returned first. The entire cache line is eventually refilled from memory.

3.5.6. Read and Write Responses

For any Avalon-MM agent, commands must be processed in a hazard-free manner. Read and write responses issue in the order in which commands they were accepted.

3.5.6.1. Transaction Order for Avalon-MM Read and Write Responses (Hosts and Agents)

For any Avalon-MM host:

- The *Avalon Interface Specifications* guarantees that commands to the same agent reach the agent in command issue order, and the agent responds in command issue order.

- Different agents may receive and respond to commands in a different order than which the host issues them. When successful, the agent responds in command issue order.

- Responses (if present) return in command issue order, regardless of whether the read or write commands are for the same or different agents.

- The *Avalon Interface Specifications* does not guarantee transaction order between different hosts.

3.5.6.2. Avalon-MM Read and Write Responses Timing Diagram

The following diagram shows command acceptance and command issue order for Avalon-MM read and write responses. Because the read and write interfaces share the `response` signal, an interface cannot issue or accept a write response and a read response in the same clock cycle.

Read responses, send one response for each `readdata`. A read burst length of *<N>* results in *<N>* responses.

intel.

Write responses, send one response for each write command. A write burst results in only one response. The agent interface sends the response after accepting the final write transfer in the burst. When an interface includes the `writeresponsevalid` signal, all write commands must complete with write responses.

Figure 16. Avalon-MM Read and Write Responses Timing Diagram
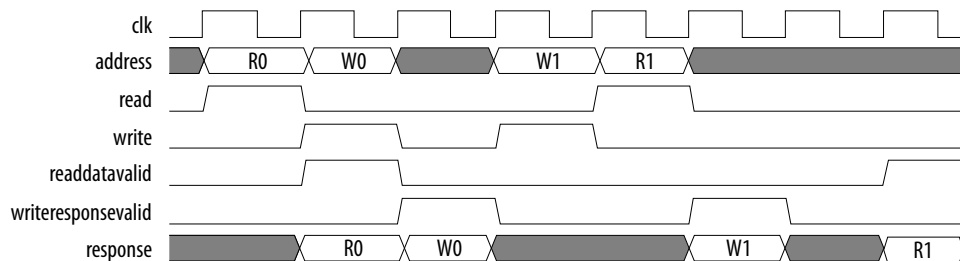
3.5.6.2.1. minimumResponseLatency Timing Diagram with readdatavalid or writeresponsevalid

For interfaces with `readdatavalid` or `writeresponsevalid`, the default a one-cycle `minimumResponseLatency` can lead to difficulty closing timing on Avalon-MM hosts.

The following timing diagrams show the behavior for a `minimumResponseLatency` of 1 or 2 cycles. Note that the actual response latency can also be greater than the minimum allowed value as these timing diagrams illustrate.

**Figure 17. minimumResponseLatency Equals One Cycle**

Figure 18. minimumResponseLatency Equals Two Cycles

Compatibility

Interfaces with the same `minimumResponseLatency` are interoperable without any adaptation. If the host has a higher `minimumResponseLatency` than the agent, use pipeline registers to compensate for the differences. The pipeline registers should

delay `readdata` from the agent. If the agent has a higher
`minimumResponseLatency` than the host, the interfaces are interoperable without
adaptation.

3.6. Address Alignment

The interconnect only supports aligned accesses. A host can only issue addresses that
are a multiple of its data width in symbols. A host can write partial words by
deasserting some `byteenables`. For example, the `byteenables` of a write of 2
bytes at address 2 is `4'b1100`.

3.7. Avalon-MM Agent Addressing

Dynamic bus sizing manages data during transfers between host-agent pairs of
differing data widths. Agent data are aligned in contiguous bytes in the host address
space.

If the host data width is wider than the agent data width, words in the host address
space map to multiple locations in the agent address space. For example, a 32-bit
host read from a 16-bit agent results in two read transfers on the agent side. The
reads are to consecutive addresses.

If the host is narrower than the agent, then the interconnect manages the agent byte
lanes. During host read transfers, the interconnect presents only the appropriate byte
lanes of agent data to the narrower host. During host write transfers, the interconnect
automatically asserts the `byteenable` signals to write data only to the specified
agent byte lanes.

Agents must have a data width of 8, 16, 32, 64, 128, 256, 512 or 1024 bits. The
following table shows the alignment for agent data of various widths within a 32-bit
host performing full-word accesses. In this table, `OFFSET[N]` refers to a agent word
size offset into the agent address space.

Table 12. **Dynamic Bus Sizing Host-to-Agent Address Mapping**

Host Byte Address (1)	Access	32-Bit Host Data		
		When Accessing an 8-Bit Agent Interface	When Accessing a 16-Bit Agent Interface	When Accessing a 64-Bit Agent Interface
0x00	1	$\text{OFFSET[0]}_{7..0}$	$\text{OFFSET[0]}_{15..0}$ **(2)**	$\text{OFFSET[0]}_{31..0}$
	2	$\text{OFFSET[1]}_{7..0}$	$\text{OFFSET[1]}_{15..0}$	—
	3	$\text{OFFSET[2]}_{7..0}$	—	—
	4	$\text{OFFSET[3]}_{7..0}$	—	—
0x04	1	$\text{OFFSET[4]}_{7..0}$	$\text{OFFSET[2]}_{15..0}$	$\text{OFFSET[0]}_{63..32}$
	2	$\text{OFFSET[5]}_{7..0}$	$\text{OFFSET[3]}_{15..0}$	—
	3	$\text{OFFSET[6]}_{7..0}$	—	—
	4	$\text{OFFSET[7]}_{7..0}$	—	—
0x08	1	$\text{OFFSET[8]}_{7..0}$	$\text{OFFSET[4]}_{15..0}$	$\text{OFFSET[1]}_{31..0}$
	2	$\text{OFFSET[9]}_{7..0}$	$\text{OFFSET[5]}_{15..0}$	—

continued...

Send Feedback

Host Byte Address (1)	Access	32-Bit Host Data		
		When Accessing an 8-Bit Agent Interface	**When Accessing a 16-Bit Agent Interface**	**When Accessing a 64-Bit Agent Interface**
	3	OFFSET[10]$_{7..0}$	—	—
	4	OFFSET[11]$_{7..0}$	—	—
0x0C	1	OFFSET[12]$_{7..0}$	OFFSET[6]$_{15..0}$	OFFSET[1]$_{63..32}$
	2	OFFSET[13]$_{7..0}$	OFFSET[7]$_{15..0}$	—
	3	OFFSET[14]$_{7..0}$	—	—
	4	OFFSET[15]$_{7..0}$	—	—
And so on		And so on	And so on	And so on

Notes:
1. Although the host issues byte addresses, the host accesses full 32-bit words.
2. For all agent entries, [*<n>*] is the word offset and the subscript values are the bits in the word.

intel.

4. Avalon Interrupt Interfaces

Avalon Interrupt interfaces allow agent components to signal events to host components. For example, a DMA controller can interrupt a processor after completing a DMA transfer.

4.1. Interrupt Sender

An interrupt sender drives a single interrupt signal to an interrupt receiver. The timing of the `irq` signal must be synchronous to the rising edge of its associated clock. `irq` has no relationship to any transfer on any other interface. `irq` must be asserted until acknowledged on the associated Avalon-MM agent interface.

Interrupts are component specific. The receiver typically determines the appropriate response by reading an interrupt status register from an Avalon-MM agent interface.

4.1.1. Avalon Interrupt Sender Signal Roles

Table 13.   Interrupt Sender Signal Roles

Signal Role	Width	Direction	Required	Description
`irq` `irq_n`	1-32	Output	Yes	Interrupt Request. An interrupt sender drives an interrupt signal to an interrupt receiver.

4.1.2. Interrupt Sender Properties

Table 14.   Interrupt Sender Properties

Property Name	Default Value	Legal Values	Description
`associatedAddressablePoint`	N/A	Name of Avalon-MM agent on this component.	The name of the Avalon-MM agent interface that provides access to the registers to service the interrupt.
`associatedClock`	N/A	Name of a clock interface on this component.	The name of the clock interface to which this interrupt sender is synchronous. The sender and receiver may have different values for this property.
`associatedReset`	N/A	Name of a reset interface on this component.	The name of the reset interface to which this interrupt sender is synchronous.

![intel]

4.2. Interrupt Receiver

An interrupt receiver interface receives interrupts from interrupt sender interfaces. Components with Avalon-MM host interfaces can include an interrupt receiver to detect interrupts asserted by agent components with interrupt sender interfaces. The interrupt receiver accepts interrupt requests from each interrupt sender as a separate bit.

4.2.1. Avalon Interrupt Receiver Signal Roles

Table 15.   Interrupt Receiver Signal Roles

Signal Role	Width	Direction	Required	Description
irq	1–32	Input	Yes	irq is an *<n>*-bit vector, where each bit corresponds directly to one IRQ sender with no inherent assumption of priority.

4.2.2. Interrupt Receiver Properties

**Table 16.   Interrupt Receiver Properties**

Property Name	Default Value	Legal Values	Description
associatedAddressable Point	N/A	Name of Avalon-MM host interface	The name of the Avalon-MM host interface used to service interrupts received on this interface.
associatedClock	N/A	Name of an Avalon Clock interface	The name of the Avalon Clock interface to which this interrupt receiver is synchronous. The sender and receiver may have different values for this property.
associatedReset	N/A	Name of an Avalon Reset interface	The name of the reset interface to which this interrupt receiver is synchronous.

4.2.3. Interrupt Timing

The Avalon-MM host services the priority 0 interrupt before the priority 1 interrupt.

Figure 19.   Interrupt Timing

In the following figure, interrupt 0 has higher priority. The interrupt receiver is in the process of handling int1 when int0 is asserted. The int0 handler is called and completes. Then, the int1 handler resumes. The diagram shows int0 deasserts at time 1. int1 deasserts at time 2.

intel.

5. Avalon Streaming Interfaces

You can use Avalon Streaming (Avalon-ST) interfaces for components that drive high-bandwidth, low-latency, unidirectional data. Typical applications include multiplexed streams, packets, and DSP data. The Avalon-ST interface signals can describe traditional streaming interfaces supporting a single stream of data without knowledge of channels or packet boundaries. The interface can also support more complex protocols capable of burst and packet transfers with packets interleaved across multiple channels.

Note: If you need a high-performance data streaming interface, refer to Chapter 6 *Avalon Streaming Credit Interfaces*.

Figure 20.   Avalon-ST Interface - Typical Application of the Avalon-ST Interface
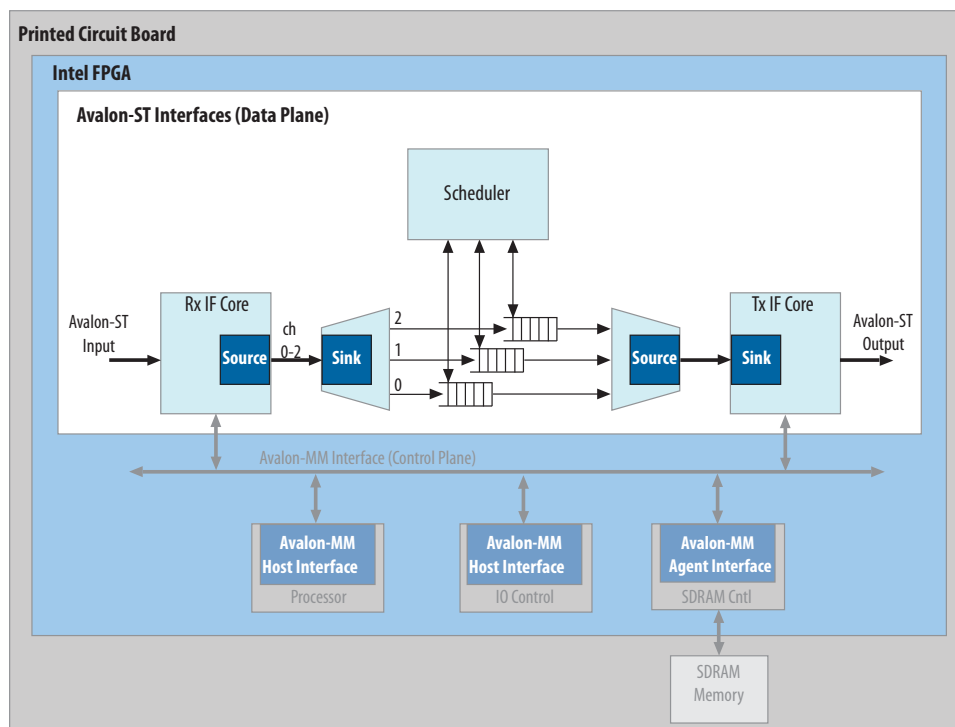
All Avalon-ST source and sink interfaces are not necessarily interoperable. However, if two interfaces provide compatible functions for the same application space, adapters are available to allow them to interoperate.

Avalon-ST interfaces support datapaths requiring the following features:

- Low-latency, high-throughput point-to-point data transfer
- Multiple channels support with flexible packet interleaving
- Sideband signaling of channel, error, and start and end of packet delineation
- Support for data bursting
- Automatic interface adaptation

5.1. Terms and Concepts

The Avalon-ST interface protocol defines the following terms and concepts:

- Avalon Streaming System—An Avalon Streaming system contains one or more Avalon-ST connections that transfer data from a source interface to a sink interface. The system shown above consists of Avalon-ST interfaces to transfer data from the system input to output. Avalon-MM control and status register interfaces provide for software control.

- Avalon Streaming Components—A typical system using Avalon-ST interfaces combines multiple functional modules, called components. The system designer configures the components and connects them together to implement a system.

- Source and Sink Interfaces and Connections—When two components connect, the data flows from the source interface to the sink interface. The *Avalon Interface Specifications* calls the combination of a source interface connecting to a sink interface a *connection*.

- Backpressure—Backpressure allows a sink to signal a source to stop sending data. Support for backpressure is optional. The sink uses backpressure to stop the flow of data for the following reasons:
  - When the sink FIFOs are full
  - When there is congestion on its output interface

- Transfers and Ready Cycles—A transfer results in data and control propagation from a source interface to a sink interface. For data interfaces, a ready cycle is a cycle during which the sink can accept a transfer.

- Symbol—A symbol is the smallest unit of data. For most packet interfaces, a symbol is a byte. One or more symbols make up the single unit of data transferred in a cycle.

- Channel—A channel is a physical or logical path or link through which information passes between two ports.

- Beat—A beat is a single cycle transfer between a source and sink interface made up of one or more symbols.

- Packet—A packet is an aggregation of data and control signals that a source transmits simultaneously. A packet may contain a header to help routers and other network devices direct the packet to the correct destination. The application defines the packet format, not this specification. Avalon-ST packets can be variable in length and can be interleaved across a connection. With an Avalon-ST interfaces, the use of packets is optional.

5.2. Avalon Streaming Interface Signal Roles

Each signal in an Avalon streaming source or sink interface corresponds to one Avalon streaming signal role. An Avalon streaming interface may contain only one instance of each signal role. All Avalon streaming signal roles apply to both sources and sinks and have the same meaning for both.

Table 17.   Avalon Streaming Interface Signals

In the following table, all signal roles are active high.

Signal Role	Width	Direction	Required	Description
Fundamental Signals				
channel	1 – 128	Source → Sink	No	The channel number for data being transferred on the current cycle. If an interface supports the channel signal, the interface must also define the maxChannel parameter.
data	1 – 8,192	Source → Sink	No	The data signal from the source to the sink, typically carries the bulk of the information being transferred. Parameters further define the contents and format of the data signal.
error	1 – 256	Source → Sink	No	A bit mask to mark errors affecting the data being transferred in the current cycle. A single bit of the error signal masks each of the errors the component recognizes. The errorDescriptor defines the error signal properties.
ready	1	Sink → Source	No	Asserts high to indicate that the sink can accept data. ready is asserted by the sink on cycle <n> to mark cycle <n + readyLatency> as a ready cycle. The source may only assert valid and transfer data during ready cycles. Sources without a ready input do not support backpressure. Sinks without a ready output never need to backpressure.
valid	1	Source → Sink	No	The source asserts this signal to qualify all other source to sink signals. The sink samples data and other source-to-sink signals on ready cycles where valid is asserted. All other cycles are ignored. Sources without a valid output implicitly provide valid data on every cycle that a sink is not asserting backpressure. Sinks without a valid input expect valid data on every cycle that they are not backpressuring.
Packet Transfer Signals				
empty	1 – 10	Source → Sink	No	Indicates the number of symbols that are empty, that is, do not represent valid data. The empty signal is not necessary on interfaces where there is one symbol per beat.
endofpacket	1	Source → Sink	No	Asserted by the source to mark the end of a packet.
startofpacket	1	Source → Sink	No	Asserted by the source to mark the beginning of a packet.

Send Feedback

5.3. Signal Sequencing and Timing

5.3.1. Synchronous Interface

All transfers of an Avalon-ST connection occur synchronous to the rising edge of the associated clock signal. All outputs from a source interface to a sink interface, including the `data`, `channel`, and `error` signals, must be registered on the rising edge of clock. Inputs to a sink interface do not have to be registered. Registering signals at the source facilitates high frequency operation.

5.3.2. Clock Enables

Avalon-ST components typically do not include a clock enable input. The Avalon-ST signaling itself is sufficient to determine the cycles that a component should and should not be enabled. Avalon-ST compliant components may have a clock enable input for their internal logic. However, components using clock enables must ensure that the timing of the interface adheres to the protocol.

5.4. Avalon-ST Interface Properties

Table 18.   Avalon-ST Interface Properties

Property Name	Default Value	Legal Values	Description
associatedClock	1	Clock interface	The name of the Avalon Clock interface to which this Avalon-ST interface is synchronous.
associatedReset	1	Reset interface	The name of the Avalon Reset interface to which this Avalon-ST interface is synchronous.
beatsPerCycle	1	1,2,4,8	Specifies the number of beats transferred in a single cycle. This property allows you to transfer 2 separate, but correlated streams using the same `start_of_packet`, `end_of_packet`, `ready` and `valid` signals. `beatsPerCycle` is a rarely used feature of the Avalon-ST protocol.
dataBitsPerSymbol	8	1 – 512	Defines the number of bits per symbol. For example, byte-oriented interfaces have 8-bit symbols. This value is not restricted to be a power of 2.
emptyWithinPacket	false	true, false	When true, `empty` is valid for the entire packet.
errorDescriptor	0	List of strings	A list of words that describe the error associated with each bit of the error signal. The length of the list must be the same as the number of bits in the error signal. The first word in the list applies to the highest order bit. For example, "`crc, overflow`" means that bit[1] of `error` indicates a CRC error. Bit[0] indicates an overflow error.
firstSymbolInHigh OrderBits	true	true, false	When true, the first-order symbol is driven to the most significant bits of the data interface. The highest-order symbol is labeled `D0` in this specification. When this property is set to false, the first symbol appears on the low bits. D0 appears at `data[7:0]`. For a 32-bit bus, if true, `D0` appears on bits[31:24].

continued...

Property Name	Default Value	Legal Values	Description
maxChannel	0	0 – 255	Maximum number of channels that a data interface can support.
readyLatency	0	0 – 8	Defines the relationship between the assertion of a `ready` signal and the assertion of a `valid` signal. If `readyLatency = <n>` where n > 0, `valid` can be asserted only <n> cycles after assertion of `ready`. For example, if `readyLatency = 1`, when the sink asserts `ready`, the source needs to respond with a `valid` assertion at least 1 cycle after it sees the `ready` assertion from the sink.
readyAllowance[1]	0	0 – 8	Defines the number of transfers that the sink can capture after `ready` is deasserted. When `readyAllowance = 0`, the sink cannot accept any transfers after `ready` is deasserted. If `readyAllowance = <n>` where <n> is greater than 0, the sink can accept up to <n> transfers after `ready` is deasserted.

Note: If you generate an Avalon streaming interconnect with Avalon streaming source/sink BFMs or custom components and these BFMs or custom components have different readyLatency requirements, Platform Designer will insert adapters in the generated interconnect to accommodate the readyLatency difference between the source and sink interfaces. It is expected that your source and sink logic adheres to the properties of the generated interconnect.

5.5. Typical Data Transfers

This section defines the transfer of data from a source interface to a sink interface. In all cases, the data source and the data sink must comply with the specification. The data sink is not responsible for detecting source protocol errors.

5.6. Signal Details

The figure shows the signals that Avalon-ST interfaces typically includes. A typical Avalon-ST source interface drives the `valid`, `data`, `error`, and `channel` signals to the sink. The sink can apply backpressure with the `ready` signal.

---

[1] • If `readyLatency = 0`, `readyAllowance` can be 0 or greater than 0.
   • If `readyLatency > 0`, `readyAllowance` must be equal to or greater than `readyLatency`.
   • If the source or the sink do not specify a value for `readyAllowance` then `readyAllowance = readyLatency`. Designs do not require the addition of `readyAllowance` unless you want the source or the sink to take advantage of this feature.

Figure 21.   Typical Avalon-ST Interface Signals


More details about these signals:

- `ready`—On interfaces supporting backpressure, the sink asserts `ready` to mark the cycles where transfers may take place. If `ready` is asserted on cycle <n>, cycle <n + readyLatency> is considered a ready cycle.

- `valid`—The `valid` signal qualifies valid data on any cycle with data transferring from source to sink. On each valid cycle the sink samples the `data` signal and other source to sink signals.

- `data`—The `data` signal carries the bulk of the information transferred from the source to the sink. The data signal consists of one or more symbols transferred on every clock cycle. The `dataBitsPerSymbol` parameter defines how the data signal is divided into symbols.

- `error`—In the `error` signal, each bit corresponds to a possible error condition. A value of 0 on any cycle indicates error-free data on that cycle. This specification does not define the action that a component takes when an error is detected.

- `channel`—The source drives the optional `channel` signal to indicate to which channel the data belongs. The meaning of `channel` for a given interface depends on the application. In some applications, `channel` indicates the interface number. In other applications, `channel` indicates the page number or timeslot. When the `channel` signal is used, all the data transferred in each active cycle belongs to the same channel. The source may change to a different channel on successive active cycles.

  Interfaces that use the `channel` signal must define the `maxChannel` parameter to indicate the maximum channel number. If the number of channels an interface supports changes dynamically, `maxChannel` indicates the maximum number the interface can support.

5.7. Data Layout

Figure 22.   Data Symbols

The following figure shows a 64-bit data signal with `dataBitsPerSymbol`=16. Symbol 0 is the most significant symbol.

The Avalon Streaming interface supports both the big-endian and little-endian modes. The figure below is an example of the big-endian mode, where Symbol 0 is in the high-order bits.

**Figure 23.** **Layout of Data**

The timing diagram in the following figure shows a 32-bit example where `dataBitsPerSymbol=8`, and `beatsPerCycle=1`.
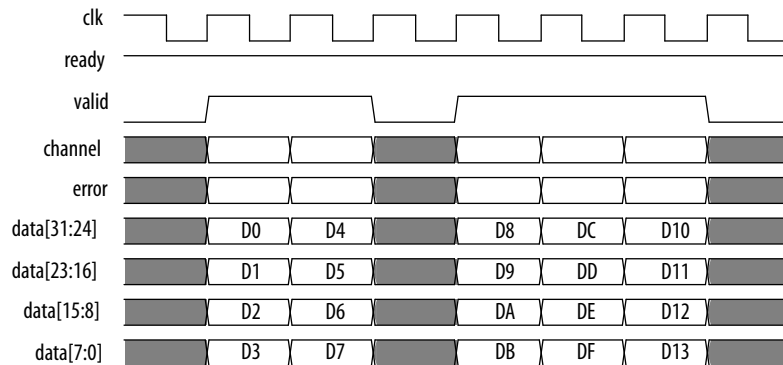
5.8. Data Transfer without Backpressure

The data transfer without backpressure is the most basic of Avalon-ST data transfers. On any given clock cycle, the source interface drives the `data` and the optional `channel` and `error` signals, and asserts `valid`. The sink interface samples these signals on the rising edge of the reference clock if `valid` is asserted.

Figure 24. **Data Transfer without Backpressure**

5.9. Data Transfer with Backpressure

The sink asserts `ready` for a single clock cycle to indicate it is ready for an active cycle. If the sink is ready for data, the cycle is a ready cycle. During a ready cycle, the source may assert `valid` and provide data to the sink. If the source has no data to send, the source deasserts `valid` and can drive `data` to any value.

Interfaces that support backpressure define the `readyLatency` parameter to indicate the number of cycles from the time that ready is asserted until valid data can be driven. If the `readyLatency` is nonzero, cycle <n + readyLatency> is a ready cycle if `ready` is asserted on cycle <n>.

When `readyLatency = 0`, data transfer only happens when `ready` and `valid` are asserted on the same cycle. In this mode, the source does not receive the sink's ready signal before sending valid data. The source provides the data and asserts `valid` whenever the source has valid data. The source waits for the sink to capture the data and assert `ready`. The source can change the data at any time. The sink only captures input data from the source when `ready` and `valid` are both asserted.

When `readyLatency >= 1`, the sink asserts `ready` before the `ready` cycle itself. The source can respond during the appropriate subsequent cycle by asserting `valid`. The source may not assert `valid` during cycles that are not `ready` cycles.

`readyAllowance` defines the number of transfers that the sink can capture when `ready` is deasserted. When `readyAllowance = 0`, the sink cannot accept any transfers after `ready` is deasserted. If `readyAllowance = <n>` where n > 0, the sink can accept up to <n> transfers after `ready` is deasserted.

5.9.1. Data Transfers Using readyLatency and readyAllowance

The following rules apply when transferring data with `readyLatency` and `readyAllowance`.

- If `readyLatency` is 0, `readyAllowance` can be greater than or equal to 0.

- If `readyLatency` is greater than 0, `readyAllowance` can be greater than or equal to `readyLatency`.

When `readyLatency = 0` and `readyAllowance = 0`, data transfers occur only when both `ready` and `valid` are asserted. In this case, the source does not receive the sink's `ready` signal before sending valid data. The source provides the data and asserts `valid` whenever possible. The source waits for the sink to capture the data and assert `ready`. The source can change the data at any time. The sink only captures input data from the source when `ready` and `valid` are both asserted.

Figure 25.   readyLatency = 0, readyAllowance = 0

When `readyLatency = 0` and `readyAllowance = 0` the source can assert `valid` at any time. The sink captures the data from source only when `ready = 1`.

The following figure demonstrates these events:

1. In cycle 1 the source provides data and asserts `valid`.

2. In cycle 2, the sink asserts `ready` and `D0` transfers.

3. In cycle 3, `D1` transfers.

4. In cycle 4, the sink asserts `ready`, but the source does not drive valid data.

5. The source provides data and asserts `valid` on cycle 6.

6. In cycle 8, the sink asserts `ready`, so `D2` transfers.

7. `D3` transfers at cycle 9 and `D4` transfers at cycle 10.

Figure 26.   readyLatency = 0, readyAllowance = 1

When `readyLatency = 0` and `readyAllowance = 1` the sink can capture one more data transfer after `ready = 0`.

The following figure demonstrates these events:

1. In cycle 1 the source provides data and asserts `valid` while the sink asserts `ready`. `D0` transfers.

2. `D1` is transferred in cycle 2.

3. In cycle 3, `ready` deasserts, however since `readyAllowance = 1` one more transfer is allowed, so `D2` transfers.

4. In cycle 5 both `valid` and `ready` assert, so `D3` transfers.

5. In cycle 6, the source deasserts `valid`, so no data transfers.

6. In cycle 7, `valid` asserts and `ready` deasserts, however since `readyAllowance = 1` one more transfer is allowed, so `D4` transfers.


Figure 27.   readyLatency = 1, readyAllowance = 2

When `readyLatency = 1` and `readyAllowance = 2` the sink can transfer data one cycle after `ready` asserts, and two more cycles of transfers are allowed after `ready` deasserts.

The following figure demonstrates these events:

1. In cycle 0 the sink asserts `ready`.

2. In cycle 1, the source provides data and asserts `valid`. The transfer occurs immediately.

3. In cycle 3, the sink deasserts `ready`, but the source is still asserting `valid`, and drives valid data because the sink can capture data two cycles after `ready` deasserts.

4. In cycle 6, the sink asserts `ready`.

5. In cycle 7, the source provides data and asserts `valid`. This data is accepted.

6. In cycle 10, the sink has deasserted `ready`, but the source asserts `valid` and drives valid data because the sink can capture data two cycles after `ready` deasserts.

Adaptation Requirements

The following table describes whether source and sink interfaces require adaptation.

Table 19.   Source/Sink Adaptation Requirements

readyLatency	readyAllowance	Adaptation
Source `readyLatency` = Sink `readyLatency`	Source `readyAllowance` = Sink `readyAllowance`	**No adaptation required:** The sink can capture all transfers.
	Source `readyAllowance` > Sink `readyAllowance`	**Adaptation required:** After `ready` is deasserted, the source can send more transfers than the sink can capture.
	Source `readyAllowance` < Sink `readyAllowance`	**No adaptation required:** After `ready` is deasserted, the sink can capture more transfers than the source can send.
Source `readyLatency` > Sink `readyLatency`	Source `readyAllowance` = Sink `readyAllowance`	**No adaptation required:** After `ready` is asserted, the source starts sending later than the sink can capture. After `ready` is deasserted, the source can send as many transfers as the sink can capture.
	Source `readyAllowance`> Sink `readyAllowance`	**Adaptation required:** After `ready` is deasserted, the source can send more transfers than the sink can capture.
	Source `readyAllowance`< Sink `readyAllowance`	**No adaptation required:** After `ready` is deasserted, the source sends fewer transfers than the sink can capture.
Source `readyLatency` < Sink`readyLatency`	Source `readyAllowance` = Sink `readyAllowance`	**Adaptation required:** The source can start sending transfers before sink can capture.
	Source `readyAllowance`> Sink `readyAllowance`	**Adaptation required:** The source can start sending transfers before the sink can capture. Also, after `ready` is deasserted, the source can send more transfers than the sink can capture.
	Source `readyAllowance` < Sink `readyAllowance`	**Adaptation required:** The source can start sending transfers before the sink can capture.

5.9.2. Data Transfers Using readyLatency

If the source or the sink do not specify a value for `readyAllowance` then `readyAllowance`= `readyLatency`. Designs that use source and sink do not require the addition of `readyAllowance` unless you want the source or the sink to take advantage of this feature.

Figure 28.   Transfer with Backpressure, readyLatency=0

The following figure illustrates these events:

1.   The source provides data and asserts `valid` on cycle 1, even though the sink is not ready.

2.   The source waits until cycle 2, when the sink does assert `ready`, before moving onto the next data cycle.

3.   In cycle 3, the source drives data on the same cycle and the sink is ready to receive data. The transfer occurs immediately.

4.   In cycle 4, the sink asserts `ready`, but the source does not drive valid data.


Figure 29.   Transfer with Backpressure, readyLatency=1

The following figures show data transfers with `readyLatency=1` and `readyLatency=2`, respectively. In both these cases, `ready` is asserted before the ready cycle, and the source responds 1 or 2 cycles later by providing data and asserting `valid`. When `readyLatency` is not 0, the source must deassert `valid` on non-ready cycles.


**Figure 30.   Transfer with Backpressure, readyLatency=2**

5.10. Packet Data Transfers

The packet transfer property adds support for transferring packets from a source interface to a sink interface. Three additional signals are defined to implement the packet transfer. Both the source and sink interfaces must include these additional signals to support packets. You can only connect source and sink interfaces with

intel.

matching packet properties. Platform Designer does not automatically add the
`startofpacket` , `endofpacket`, and `empty` signals to source or sink interfaces that
do not include these signals.

Figure 31.   Avalon-ST Packet Interface Signals

5.11. Signal Details

- `startofpacket`—All interfaces supporting packet transfers require the
  `startofpacket` signal. `startofpacket` marks the active cycle containing the
  start of the packet. This signal is only interpreted when `valid` is asserted.

- `endofpacket`—All interfaces supporting packet transfers require the
  `endofpacket` signal. `endofpacket` marks the active cycle containing the end of
  the packet. This signal is only interpreted when `valid` is asserted.
  `startofpacket` and `endofpacket` can be asserted in the same cycle. No idle
  cycles are required between packets. The `startofpacket` signal can follow
  immediately after the previous `endofpacket` signal.

- `empty`—The optional `empty` signal indicates the number of symbols that are
  empty during the `endofpacket` cycle. The sink only checks the value of the
  `empty` during active cycles that have `endofpacket` asserted. The empty symbols
  are always the last symbols in `data`, those carried by the low-order bits when
  `firstSymbolInHighOrderBits` = true. The `empty` signal is required on all
  packet interfaces whose `data` signal carries more than one symbol of data and
  have a variable length packet format. The size of the `empty` signal in bits is
 $ceil[log_2(<symbols\ per\ cycle>)]$.

5.12. Protocol Details

Packet data transfer follows the same protocol as the typical data transfer with the addition of the `startofpacket`, `endofpacket`, and `empty`.

Figure 32.   Packet Transfer

The following figure illustrates the transfer of a 17-byte packet from a source interface to a sink interface, where `readyLatency=0`. This timing diagram illustrates the following events:

1.   Data transfer occurs on cycles 1, 2, 4, 5, and 6, when both `ready` and `valid` are asserted.

2.   During cycle 1, `startofpacket` is asserted. The first 4 bytes of packet are transferred.

3.   During cycle 6, `endofpacket` is asserted. `empty` has a value of 3. This value indicates that this is the end of the packet and that 3 of the 4 symbols are empty. In cycle 6, the high-order byte, data[31:24] drives valid data.

intel.
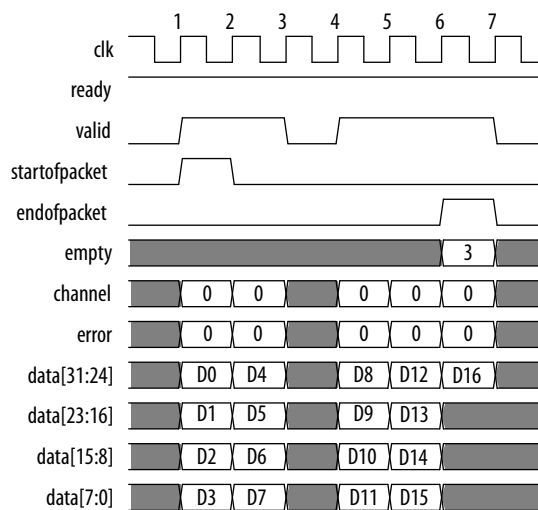
6. Avalon Streaming Credit Interfaces

Avalon Streaming Credit interfaces are for use with components that drive high-bandwidth, low-latency, unidirectional data. Typical applications include multiplexed streams, packets, and DSP data. The Avalon Streaming Credit interface signals can describe traditional streaming interfaces supporting a single stream of data, without knowledge of channels or packet boundaries. The interface can also support more complex protocols capable of burst and packet transfers with packets interleaved across multiple channels.

All Avalon Streaming Credit source and sink interfaces are not necessarily interoperable. However, if two interfaces provide compatible functions for the same application space, adapters are available to allow them to interoperate.

You can also connect the Avalon Streaming Credit source to an Avalon Streaming sink via an adapter. Similarly, you can connect an Avalon Streaming source to an Avalon Streaming Credit sink via an adapter.

Avalon Streaming Credit interfaces support datapaths requiring the following features:

- Low-latency, high-throughput point-to-point data transfer
- Multiple channels support with flexible packet interleaving
- Sideband signaling of channel, error, and start and end of packet delineation
- Support for data bursting
- User signals as sideband signals for functionality users define

6.1. Terms and Concepts

The Avalon Streaming Credit interface protocol defines the following terms and concepts:

- **Avalon Streaming Credit System**— An Avalon Streaming Credit system contains one or more Avalon Streaming Credit connections that transfer data from a source interface to a sink interface.

- **Avalon Streaming Credit Components**— A typical system using Avalon Streaming interfaces combines multiple functional modules, called components. The system designer configures the components and connects them together to implement a system.

- **Source and Sink Interfaces and Connections**—When two components are connected, credits flow from the sink to the source; and the data flows from the source interface to the sink interface. The combination of a source interface connected to a sink interface is referred to as a connection.

- **Transfers**— A transfer results in data and control propagation from a source interface to a sink interface. For data interfaces, source can start data transfer only if it has credits available. Similarly, sink can accept data only if it has outstanding credits.

ISO 9001:2015 Registered

- **Symbol**—A symbol is the smallest unit of data. One or more symbols make up the single unit of data transferred in a cycle.

- **Beat**—A beat is a single cycle transfer between a source and sink interface made up of one or more symbols.

- **Packet**—A packet is an aggregation of data and control signals that is transmitted together. A packet may contain a header to help routers and other network devices direct the packet to the correct destination. The packet format is defined by the application, not this specification. Avalon Streaming packets can be variable in length and can be interleaved across a connection. With an Avalon Streaming Credit interface, the use of packets is optional.

6.2. Avalon Streaming Credit Interface Signal Roles

Each signal in an Avalon Streaming Credit source or sink interface corresponds to one Avalon Streaming Credit signal role. An Avalon Streaming Credit interface may contain only one instance of each signal role. All Avalon Streaming Credit signal roles apply to both sources and sinks and have the same meaning for both.

Table 20.   Avalon Streaming Credit Interface Signals

Signal Name	Direction	Width	Optional / Required	Description
`update`	Sink to source	1	Required	Sink sends `update` and source updates the available credit counter. Sink sends `update` to source when a transaction is popped from its buffer. Credit counter in source is increased by the value on the credit bus from sink to source.
`credit`	Sink to source	1-9	Required	Indicates additional credit available at sink when update is asserted. This bus carries a value as specified by the sink. Width of the `credit` bus is $ceilog_2(MAX\_CREDIT + 1)$. Sink sends available credit value on this bus which indicates the number of transactions it can accept. Source captures **credit** value only if **update** signal is asserted.
`return_credit`	Source to sink	1	Required	Asserted by source to return 1 credit back to sink. *Note:* For more details, refer to Section 6.2.3 *Returning the Credits*.
`data`	Source to sink	1-16368	Required	Data is divided into symbols as per existing Avalon Streaming definition.
`valid`	Source to sink	1	Required	Asserted by the source to qualify all other source to sink signals. Source can assert `valid` only when the credit available to it is greater than 0.
`error`	Source to sink	1-256	Optional	A bit mask used to mark errors affecting the data being transferred in the current cycle. A single bit in error is used for each of the errors recognized by the component, as defined by the `errorDescriptor` property.

continued...

Signal Name	Direction	Width	Optional / Required	Description
`channel`	Source to sink	1-128	Optional	The channel number for data being transferred on the current cycle. If an interface supports the `channel` signal, it must also define the `maxChannel` parameter.
Packet Transfer Signals				
`startofpacket`	Source to sink	1	Optional	Asserted by the source to mark the start of a packet.
`endofpacket`	Source to sink	1	Optional	Asserted by the source to mark the end of a packet.
`empty`	Source to sink	ceil(log2(NUM_SYMBOLS))	Optional	Indicates the number of symbols that are empty, that is, do not represent valid data. The `empty` signal is not used on interfaces where there is one symbol per beat.
User Signals				
`<Per-Packet User Signals>`	Source to sink	1-16368	Optional	Any number of per-packet user signals can be present on source and sink interfaces. Source sets value of this signal when `startofpacket` is asserted. Source should not change the value of this signal until start of new packet. More details are in the User Signal section.
`<Per-Symbol User Signals>`	Source to sink	1-16368	Optional	Any number of per-symbol user signals can be present on source and sink. More details are in the User Signal section.

6.2.1. Synchronous Interface

All transfers of an Avalon Streaming connection occur synchronous to the rising edge of the associated clock signal. All outputs from a source interface to a sink interface, including the `data`, `channel`, and `error` signals, must be registered on the rising edge of clock. Inputs to a sink interface do not have to be registered. Registering signals at the source facilitates high-frequency operation.

Table 21.   Avalon Streaming Credit Interface Properties

| Property Name	Default Value	Legal Value	Description
`associatedClock`	1	Clock interface	The name of the Avalon Clock interface to which this Avalon Streaming interface is synchronous.
`associatedReset`	1	Reset interface	The name of the Avalon Reset interface to which this Avalon Streaming interface is synchronous.
`dataBitsPerSymbol`	8	1 – 16368	Defines the number of bits per symbol. For example, byte-oriented interfaces have 8-bit symbols. This value is not restricted to be a power of 2.
`symbolsPerBeat`	1	1 – 16368	The number of symbols that are transferred on every valid cycle.
`maxCredit`	256	1-256	The maximum number of credits that a data interface can support.

continued...

Property Name	Default Value	Legal Value	Description
errorDescriptor	0	List of strings	A list of words that describe the error associated with each bit of the error signal. The length of the list must be the same as the number of bits in the error signal. The first word in the list applies to the highest order bit. For example, "crc, overflow" means that bit[1] of error indicates a CRC error. Bit[0] indicates an overflow error.
firstSymbolInHighOrderBits	true	true, false	When true, the first-order symbol is driven to the most significant bits of the data interface. The highest-order symbol is labeled D0 in this specification. When this property is set to false, the first symbol appears on the low bits. D0 appears at data[7:0]. For a 32-bit bus, if true, D0 appears on bits[31:24].
maxChannel	0	0	The maximum number of channels that a data interface can support.

6.2.2. Typical Data Transfers

This section defines the transfer of data from a source interface to a sink interface. In all cases, the data source and the data sink must comply with the specification. It is not the responsibility of the data sink to detect source protocol errors.

The below figure shows the signals that are typically used in an Avalon Streaming Credit interface.

Figure 33. Typical Avalon Streaming Credit Signals
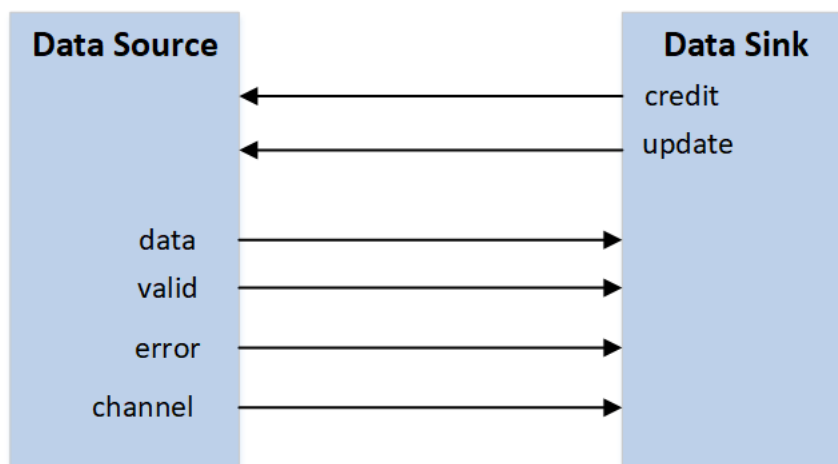
As this figure indicates, a typical Avalon Streaming Credit source interface drives the `valid`, `data`, `error`, and `channel` signals to the sink. The sink drives `update` and `credit` signals.

Figure 34. **Typical Credit and Data Transfer**

The above figure shows a typical credit and data transfer between source and sink. There can be an arbitrary delay between the sink asserting `update` and source receiving the `update`. Similarly, there can be an arbitrary delay between source asserting `valid` for data and sink receiving that data. Delay on credit path from sink to source and data path from source to sink need not be equal. These delays can be 0 cycle as well, i.e. when the sink asserts `update`, it is seen by the source in the same cycle. Conversely, when the source asserts `valid`, it is seen by the sink in the same cycle. If source has zero credits, it cannot assert `valid`. Transferred credits are cumulative. If sink has transferred credits equal to its maxCredit property, and has not received any data, it cannot assert `update` until it receives at least 1 data or has received a `return_credit` pulse from the source.

Sink cannot backpressure data from source if sink has provided credits to the source, i.e. sink must accept data from source if there are outstanding credits. Source cannot assert `valid` if it has not received any credit or exhausted the credits received, i.e. already sent the data in lieu of credits received.

If source has zero credits, source cannot start the data transfer in the same cycle it receives credits. Similarly, if sink has transferred credits equal to its maxCredit property and it receives data, sink cannot send an update in the same cycle as it received data. These restrictions have been put in place to avoid combinational loops in the implementation.

6.2.3. Returning the Credits

Avalon Streaming Credit protocol supports a `return_credit` signal. This is used by source to return the credits back to sink. Every cycle this signal is asserted, it indicates source is giving back 1 credit. If source wants to return multiple credits, this signal needs to be asserted for multiple cycles. For example, if source wants to return 10 outstanding credits, it asserts `return_credit` signal for 10 cycles. Sink should account for returned credits in its internal credit maintenance counters. Credits can be returned by source at any point in time as long as it has credits greater than 0.

The below figure exemplifies source returning credits. As shown in the figure, *outstanding_credit* is an internal counter for the source. When source returns credits, this counter is decremented.

Figure 35. **Source Returning Credits**

Note: Although the diagram above shows the returning of credits when `valid` is deasserted, `return_credit` can also be asserted while `valid` is asserted. In this case, source effectively spends 2 credits: one for `valid`, and one for `return_credit`.

6.3. Avalon Streaming Credit User Signals

User signals are optional sideband signals which flow along with data. They are considered valid only when data is valid. Given that user signals do not have any defined meaning or purpose, caution must be used while using these signals. It is the responsibility of the system designer to make sure that two IPs connected to each other agree on the roles of the user signals.

Two types of user signals are being proposed: per-symbol user signals and per-packet user signals.

6.3.1. Per-Symbol User Signal

As the name suggests, the data defines a per-symbol user signal (`symbol_user`) per symbol. Each symbol in the data can have a user signal. For example, if the number of symbols in the data is 8, and `symbol_user` width is 2 bits, the total width of the `symbol_user` signal is 16 bits.

`Symbol_user` is valid only when data is valid. Source can change this signal every cycle when data is valid. Sink can disregard the value of `symbol_user` bits for empty symbols.

If a source which has this signal is connected to a sink which does not have this signal on its interface, the signal from source remains dangling in the generated interconnect.

If a source which does not have this signal is connected to a sink which has this signal on its interface, the sink's input user signal ties to 0.

If both source and sink have equal number of symbols in the data, then the user signals for both must have equal widths. Otherwise, they cannot be connected.

Send Feedback

If a wide source is connected to a narrow sink, and both have per-symbol user signals, then both must have equal bits of user signal associated with each symbol. For example, if a 16-symbol source has 2 bits of user signal associated with each symbol (for a total of 32 bits of user signal), then a 4-symbol sink must have an 8-bit wide user signal (2 bits associated with each symbol). A data format adapter can convert the 16-symbol source data to 4-symbol sink data, and 32-bit user signal to 8-bit user signal. The data format adapter maintains the association of symbols with corresponding user signal bits.

Similarly, if a narrow source is connected to a wide sink, and both have per-symbol user signals, then both must have equal bits of user signal associated with each symbol. For example, if a 4-symbol source has 2 bits of user signal associated with each symbol (for a total of 8 bits of user signal), then a 16-symbol sink must have a 32-bit wide user signal (2 bits associated with each symbol). A data format adapter can convert the 4-symbol source data to 16-symbol sink data, and 8-bit user signal to 32-bit user signal. The data format adapter maintains the association of symbols with corresponding user signal bits. If the packet is smaller than the ratio of data widths, the data format adapter sets the value of empty accordingly. Sink should disregard the value of user bits associated with empty symbols.

6.3.2. Per-Packet User Signal

In addition to `symbol_user`, per-packet user signals (`packet_user`) can also be declared on the interface. `Packet_user` can be of arbitrary width. Unlike `symbol_user`, `packet_user` must remain constant throughout the packet, i.e. its value should be set at the start of the packet and must remain the same until the end of the packet. This restriction makes the implementation of the data format adapter simpler as it eliminates the option to replicate or chop (wide source, narrow sink) or concatenate (narrow source, wide sink) `packet_user`.

If a source has `packet_user` and sink does not, the `packet_user` from source remains dangling. In such a case, the system designer must be careful and not transmit any critical control information on this signal as it is completely or partially ignored.

If a source does not have `packet_user` and the sink does, the `packet_user` to sink is tied to 0.

intel.

7. Avalon Conduit Interfaces

Avalon Conduit interfaces group an arbitrary collection of signals. You can specify any role for conduit signals. However, when you connect conduits, the roles and widths must match, and the directions must be opposite. An Avalon Conduit interface can include input, output, and bidirectional signals. A module can have multiple Avalon Conduit interfaces to provide a logical signal grouping. Conduit interfaces can declare an associated clock. When connected conduit interfaces are in different clock domains, Platform Designer generates an error message.

Note: If possible, you should use the standard Avalon-MM or Avalon-ST interfaces instead of creating an Avalon Conduit interface. Platform Designer provides validation and adaptation for these interfaces. Platform Designer cannot provide validation or adaptation for Avalon Conduit interfaces.

Conduit interfaces typically used to drive off-chip device signals, such as an SDRAM address, data and control signals.

Figure 36.   Focus on the Conduit Interface
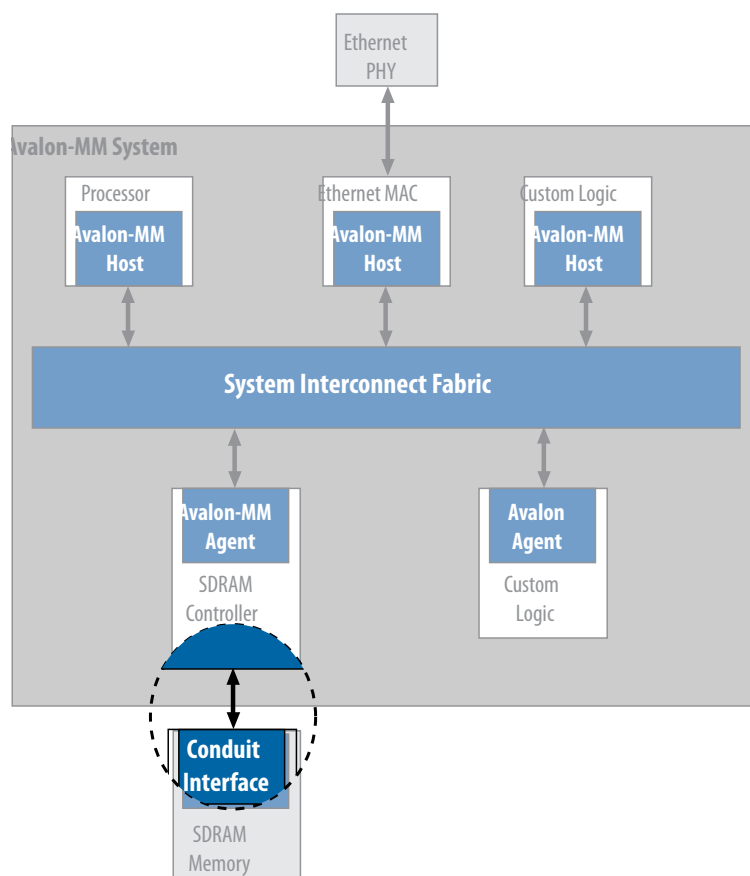
7.1. Avalon Conduit Signal Roles

Table 22.   Conduit Signal Roles

Signal Role	Width	Direction	Description
<any>	<n>	In, out, or bidirectional	A conduit interface consists of one or more input, output, or bidirectional signals of arbitrary width. Conduits can have any user-specified role. You can connect compatible Conduit interfaces inside a Platform Designer (Standard) system provided the roles and widths match and the directions are opposite.

7.2. Conduit Properties

There are no properties for conduit interfaces.

intel.

8. Avalon Tristate Conduit Interface

The Avalon Tristate Conduit Interface (Avalon-TC) is a point-to-point interface designed for on-chip controllers that drive off-chip components. This interface allows data, address, and control pins to be shared across multiple tristate devices. Sharing conserves pins in systems that have multiple external memory devices.

The Avalon-TC interface restricts the more general Avalon Conduit Interface in two ways:

- The Avalon-TC requires `request` and `grant` signals. These signals enable bus arbitration when multiple Tristate Conduit Hosts are requesting access to a shared bus.

- The pin type of a signal must be specified using suffixes appended to a signal's role. The three suffixes are: `_out`, `_in`, and `_outen`. Matching role prefixes identify signals that share the same I/O Pin. The following illustrates the naming conventions for Avalon-TC shared pins.

**Figure 37. Shared Pin Types**
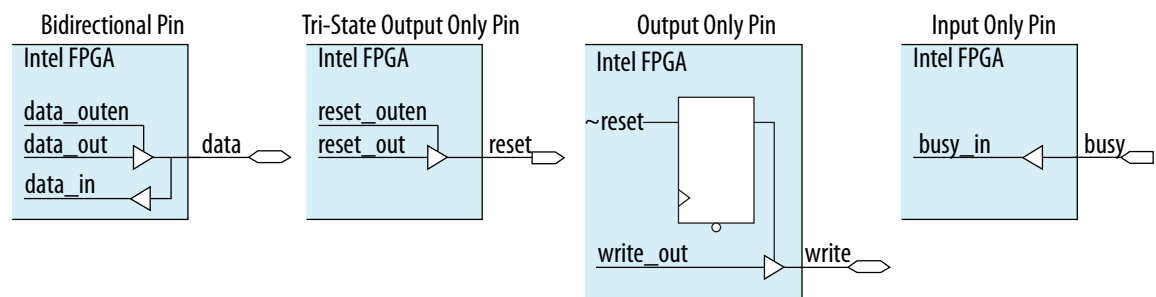
The next figure illustrates pin sharing using Avalon-TC interfaces. This figure illustrates the following points.

- The Tristate Conduit Pin Sharer includes separate Tristate Conduit Agent Interfaces for each Tristate Conduit Host. Each host and agent pair has its own `request` and `grant` signals.

- The Tristate Conduit Pin Sharer identifies signals with identical roles as tristate signals that share the same FPGA pin. In this example, the following signals are shared: `addr_out`, `data_out`, `data_in`, `read_out`, and `write_out`.

- The Tristate Conduit Pin Sharer drives a single bus including all the shared signals to the Tristate Conduit Bridge. If the widths of shared signals differ, the Tristate Conduit Pin Sharer aligns them on their 0^{th} bit. Tristate Conduit Pin Sharer drives the higher-order pins to 0 whenever the smaller signal has control of the bus.

- Signals that are not shared propagate directly through the Tristate Conduit Pin Sharer. In this example, the following signals are not shared: `chipselect0_out`, `irq0_out`, `chipselect1_out`, and `irq1_out`.

- All Avalon-TC interfaces connected to the same Tristate Conduit Pin Sharer must be in the same clock domain.

Figure 38.   Tristate Conduit Interfaces

The following illustrates the typical use of Avalon-TC Host and Agent interfaces and signal naming.
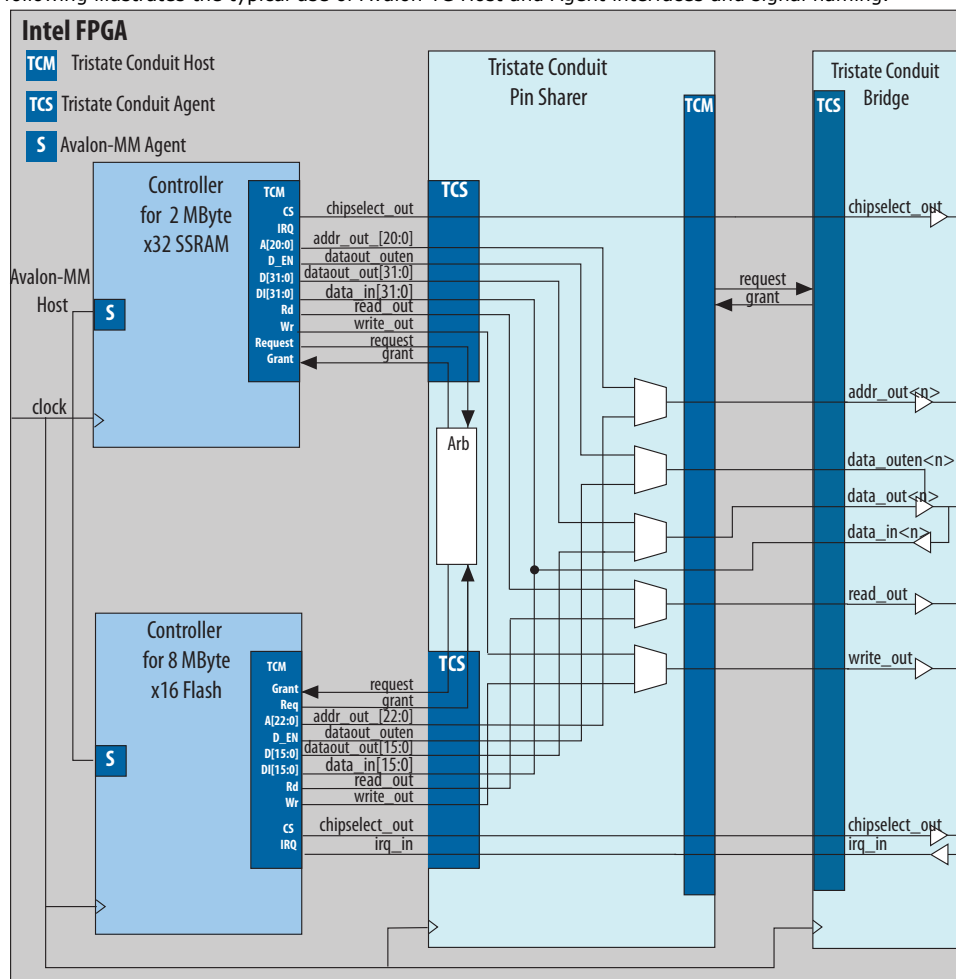
For more information about the Generic Tristate Controller and Tristate Conduit Pin Sharer, refer to the *Avalon Tristate Conduit Components User Guide*.

Related Information

Avalon Tristate Conduit Components User Guide

8.1. Avalon Tristate Conduit Signal Roles

The following table lists the signal defined for the Avalon Tristate Conduit interface. All Avalon-TC signals apply to both hosts and agents and have the same meaning for both.

Table 23.   Tristate Conduit Interface Signal Roles

Signal Role	Width	Direction	Required	Description
request	1	Host → Agent	Yes	The meaning of request depends on the state of the grant signal, as the following rules dictate. When request is asserted and grant is deasserted, request is requesting access for the current cycle. When request is asserted and grant is asserted, request is requesting access for the next cycle. Consequently, request should be deasserted on the final cycle of an access. The request signal deasserts in the last cycle of a bus access. The request signal can reassert immediately following the final cycle of a transfer. This protocol makes both rearbitration and continuous bus access possible if no other hosts are requesting access. Once asserted, request must remain asserted until granted. Consequently, the shortest bus access is 2 cycles. Refer to *Tristate Conduit Arbitration Timing* for an example of arbitration timing.
grant	1	Agent → Host	Yes	When asserted, indicates that a tristate conduit host has access to perform transactions. The grant signal asserts in response to the request signal. The grant signal remains asserted until 1 cycle following the deassertion of request.
*<name>*_in	1 – 1024	Agent → Host	No	The input signal of a logical tristate signal.
*<name>*_out	1 – 1024	Host → Agent	No	The output signal of a logical tristate signal.
*<name>*_outen	1	Host → Agent	No	The output enable for a logical tristate signal.

8.2. Tristate Conduit Properties

There are no special properties for the Avalon-TC Interface.

8.3. Tristate Conduit Timing

The following illustrates arbitration timing for the Tristate Conduit Pin Sharer. Note that a device can drive or receive valid data in the granted cycle.

Figure 39. Tristate Conduit Arbitration Timing

This figure shows the following sequence of events:

1. In cycle 1, the tristate conduit agent asserts grant. The agent drives valid data in cycles1 and 2.

2. In cycle 4, the tristate conduit agent asserts grant. The agent drives valid data in cycles 5–8.

3. In cycle 9, the tristate conduit agent asserts grant. The agent drives valid data in cycles 10–17.

4. Cycles 3, 4 and 9 do not contain valid data.

Send Feedback

intel.

A. Deprecated Signals

Deprecated signals implement functionality that is no longer required or has been superseded.

begintransfer

An output of Avalon-MM hosts. Asserted for a single cycle at the beginning of a transfer. This signal is not used and is not necessary.

chipselect or chipselect_n

`chipselect` or `chipselect_n`: The chip select signal as described below was deprecated with the release of the Avalon Tristate Conduit (Avalon-TC) interface type which includes a chip select signal.

Formerly `chipselect` was a 1-bit input to an Avalon Memory-Mapped (Avalon-MM) agent interface signaling the beginning of a read or write transfer. The current Platform Designer interconnect filters read and write signals from hosts according to the address and address map. The Platform Designer interconnect only drives read and write signals to the appropriate Avalon-MM agent, making a chip select unnecessary.

This signal dates from very early microprocessor designs. CPLDs decoded microprocessor addresses and generated chip selects for peripherals that were frequently asynchronous. With synchronous systems this signal is unnecessary.

flush

Signal removed in version 1.2 of the *Avalon Interface Specifications*. Formerly available to hosts to clear pending transfers for pipelined reads.

B. Document Revision History for the Avalon Interface Specifications

Document Version	Intel Quartus Prime Version	Changes
2021.05.27	20.1	Converted non-inclusive terms to "host" and "agent" inclusive terms for Avalon interface descriptions throughout the document.
2021.04.26	20.1	Added more clarification for the `readyLatency` property to the *Avalon-ST Interface Properties* section. Also added a Note with a description of the Avalon streaming interconnect that connects Avalon streaming source/sink BFMs or custom components to the same section.
2020.12.21	20.1	Changed references to `readyLatency` to the correct parameter `readLatency` in the *Pipelined Read Transfers with Fixed Latency* section.
2020.05.26	20.1	Added more description for the timings diagram *Figure 27* in section *Data Transfers Using* `readyLatency` *and* `readyAllowance`.
2020.05.07	20.1	Added some clarification for the timing behavior of the signal `writeresponsevalid` to the *Avalon Memory-Mapped Interface Signal Roles* section. Updated the bus widths for the `data` and `empty` signals in the *Avalon Streaming Interface Signal Roles* section.
2020.04.13	20.1	Added the chapter *Avalon Streaming Credit Interfaces*.
2020.01.03	18.1	Corrected the definition of the `burstOnBurstBoundaries` interface property. When true, the burst must begin on a multiple of the maximum burst size.
2019.10.08	18.1	Removed references to `symbolsPerBeat` because it is a deprecated parameter. Added a note in the *Data Layout* topic to clarify that the Avalon Streaming Interface supports both big-endian and little-endian modes.
2019.10.03	18.1	Corrected the property that specifies the fixed latency in the *Pipelined Read Transfers with Fixed Latency* topic. The `readyLatency` property, not the `readWaitTime` property specifies this value.
2018.09.26	18.1	In the *Write Bursts* section, added a statement saying that writes with byteenables being all 0's are passed on to the Avalon-MM agent as valid transactions.
2018.09.24	18.1	In *Avalon Memory-Mapped Interface Signal Roles*, added consecutive byte-enable support.
2018.05.22	18.0	Made the following changes: • In the *Avalon-ST Interface Properties* table, corrected the default value for `beatsPerCycle`. The default value is 1. • In the *Avalon-ST Interface Properties* table, added legal values for `beatsPerCycle`. Legal values are 1, 2, 4, and 8. • Corrected minor errors and typos.
2018.05.07	18.0	Made the following changes:

continued...

ISO 9001:2015 Registered

intel.

Document Version	Intel Quartus Prime Version	Changes
		• Added support for the `readyAllowance` parameter. • Updated the *Data Transfers with Backpressure* topic to incorporate support for the `readyAllowance` parameter. • Fixed minor errors and typos.

2018.03.22	17.1	Made the following changes:
		• Made the following changes to the *Read and Write Transfers with Waitrequest* timing diagram
		  — Removed `readdatavalid` signal which is not relevant when using `waitrequest`
		  — Moved the number 4, `readdata` and `response` forward one cycle.
		  — Aligned the `read` signal to number 1.
		• Expanded the *Transfers Using the waitrequestAllowance Property* section. Provided more complex timing diagrams.
		• Updated the discussion in the *Read Bursts* section. For reads with a `burstcount > 1`, Intel recommends asserting all `byteenables`.
		• Enhanced discussion in the *waitrequestAllowance Equals Two - Not Recommended* topic. Corrected timing diagram. Data must be held for 2 cycles starting at clock cycle 11.

November 2017	17.1	Made the following changes:
		• Updated the discussion of *Read Bursts* as follows:
		  — Qualified the statement, " When a host connects directly to a agent, a burstcount of <n>, means the agent must return <n> words of readdata to complete the burst. " This statement is true if the host connects directly to the agent. It may not be true if interconnect links the host and agent.
		  — Removed the following statement from the description of read bursts: "The byteenables presented with a read burst command apply to all cycles of the burst." This statement is no longer true. However, Intel recommends that reads with `burstcount > 1` assert all byteenables.
		• Removed the following statement form the *Pipelined Transfers* section: *Write transfers cannot be pipelined.* You can pipeline writes using the `writeresponsevalid` signal.
		• Expanded the description of read and write responses in the *Avalon-MM Read and Write Responses Timing Diagram* section.
		• Revised the description of the `reset_req` signal.
		• Changed width of `irq` from 1 bit to 1-32 bits. Both the Intel Quartus Prime Pro Edition and Intel Quartus Prime Standard Edition software support interrupt vectors.

May 2017	Quartus Prime Pro v17.1 Stratix 10 ES Editions	Made the following changes:
		• Added the following interface property parameters.
		  — `waitrequestAllowance` parameter to support high speed operation. This parameter is available for Avalon-MM interfaces. Added timing diagrams illustrating use of this parameter.
		  — `minimumResponseLatency` parameter to facilitate timing closure for Avalon-MM interface. Added timing diagrams illustrating use of this parameter.

December 2015	15.1	Made the following changes:

<div align="right">*continued...*</div>

Document Version	Intel Quartus Prime Version	Changes

		• Changed the width of the `empty` signal from a maximum of 8 bits to a maximum of 5 bits.
		• Improved the definition of the `reset_req` signal.
		• Removed the `readdatavalid` signal from the *Pipelined Read Transfer with Fixed Latency of Two Cycles* timing diagram. This signal is not relevant for fixed latency transfers.
		• Corrected equation defining the `empty` signal.
		• Made the following changes in the *Pipelined Read Transfers with Variable Latency* timing diagram:
		  — Moved the deassertion of the `read` signal to cycle 9
		  — Changed `waitrequest` to don't care in cycle 9.

March 2015	14.1	Fixed typo in Figure 1-1.

January 2015	14.1	Made the following changes:
		• Clarified address alignment example. The Avalon-MM host and agent interfaces are different widths.
		• Improved discussion of *Pipelined read Transfers with Variable Latency*. Corrected timing marker 2 which should be exactly on the rising edge of clock.
		• Improved discussion of *Pipelined Read Transfer with Fixed Latency of Two Cycles*.
		• Clarified use of `beatsPerCycle` property.
		• Corrected the address range for line-wrapped bursts. The correct address range for a 64-byte burst is 0x0–0x3C, not 0x0–0x1C.
		• Corrected description of the *Tristate Conduit Arbitration Timing* diagram in the following ways:
		  — The tristate conduit agent asserts grant, not the tristate conduit host.
		  — The final grant comes in cycle 9, not cycle 8.
		• Added a *Deprecated Signals* appendix.
		• Added read `response` signal.
		• Improved definitions of clock and reset signal types.
		• Corrected definition of clock sink properties.
		• Corrected definition of `synchronousEdges` for reset source interface.
		• Clarified the Avalon-MM `response` signal type.
		• Updated definition of `empty`. The signal must be interpreted `emptyWithinPacket` is true.
		• Edited for clarity and consistency.

June 2014	14.0	• Updated the Avalon-MM Signals table, `begintransfer`, `readdatavalid`, and `readdatavalid_n`.
		• Updated the Read and Write Transfers with Waitrequest figure:
		  — Moved deassertion of write to cycle 6.
		  — Moved assertion of `readdatavalid` and `readdata` to cycle 4.
		• Updated the Pipelined Read Transfers with Variable Latency figure:
		  — Moved assertion of `data1` to just after cycle 5, and assertion of `data2` to cycle 6.
		  — Moved assertion of `readdatavalid` to match `data1` and `data2`.

April 2014	13.01	Corrected *Read and Write Transfers with Waitrequest* In *Avalon Memory-Mapped Interfaces* chapter .
May 2013	13.0	Made the following changes:

continued...

Document Version	Intel Quartus Prime Version	Changes
		- Minor updates to *Avalon Memory-Mapped Interfaces*. - Minor updates to *Avalon Streaming Interfaces*. - Updated *Avalon Conduit Interfaces* to describe the signal roles supported by Avalon conduit interfaces. - Updated *Shared Pin Types* figure in the *Avalon Tristate Conduit Interface* chapter.
May 2011	11.0	Initial release of the *Avalon Interface Specifications*.